



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA
CAMPUS DI CESENA

What conditions enable learning? Contiguity, contingency & surprise


Cognition and Neuroscience
Academic year 2023/2024

Francesca Starita

francesca.starita2@unibo.it

Conditioned responses are anticipatory, predictive responses


development of the conditioned eyelid response



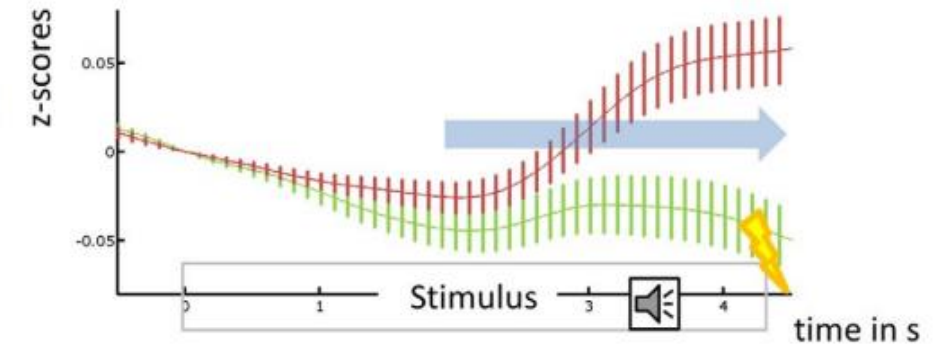
Sutton, R. S., & Barto, A. G. (2018). Reinforcement learning: An introduction. MIT press.

CS+
CS-



(a) Startle response



(b) Skin conductance response



(c) Pupil response



Leuchs, L., Schneider, M., & Spoormaker, V. I. (2019). *Psychophysiology*, 56(1), e13283. <https://doi.org/10.1111/psy.13283>

Associative learning is learning about contiguity ...

Contiguity: closeness in time between stimulus/behavior and outcome. Stimuli that are close to one another in time become associated.

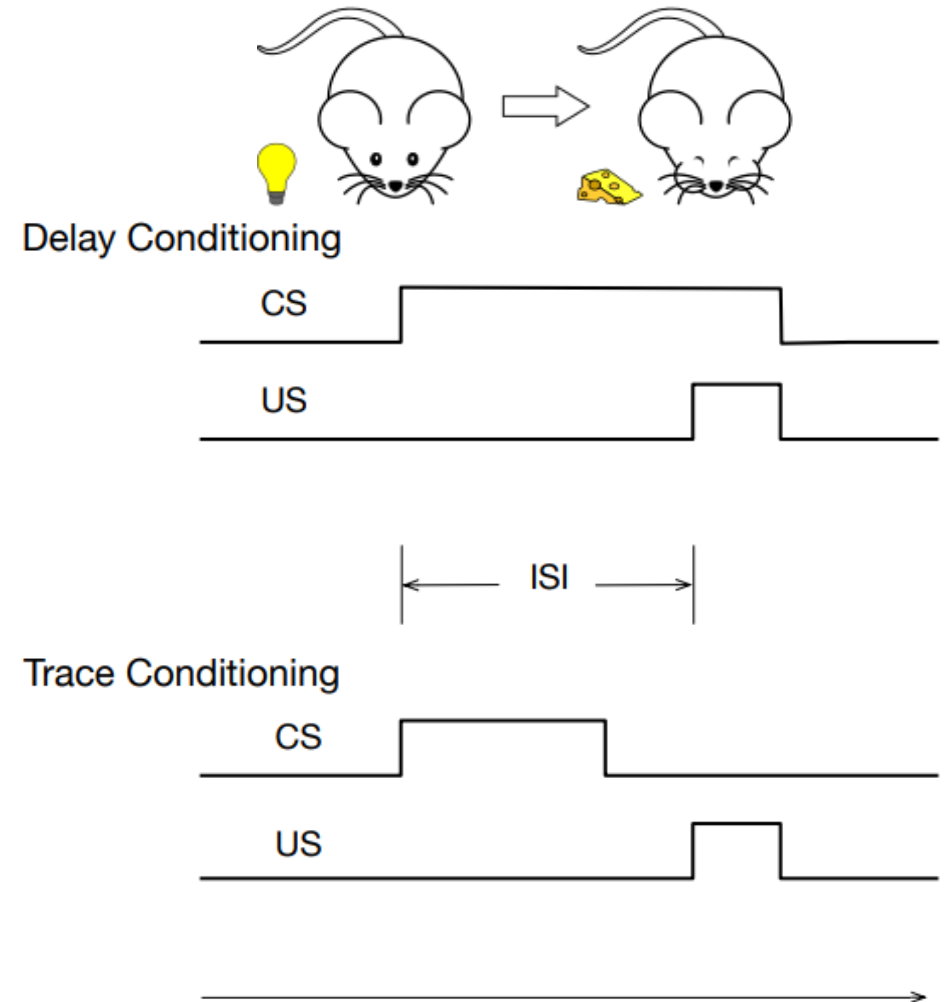
Temporal contiguity: the temporal relation between the CS and US is critical for learning

The closer in time two events occur, the more likely they will become associated.

**How would
you test
this?**

Temporal contiguity: the temporal relation between the CS and US is critical for learning


The closer in time two events occur, the more likely they will become associated.



Temporal contiguity: the temporal relation between the CS and US is critical for learning

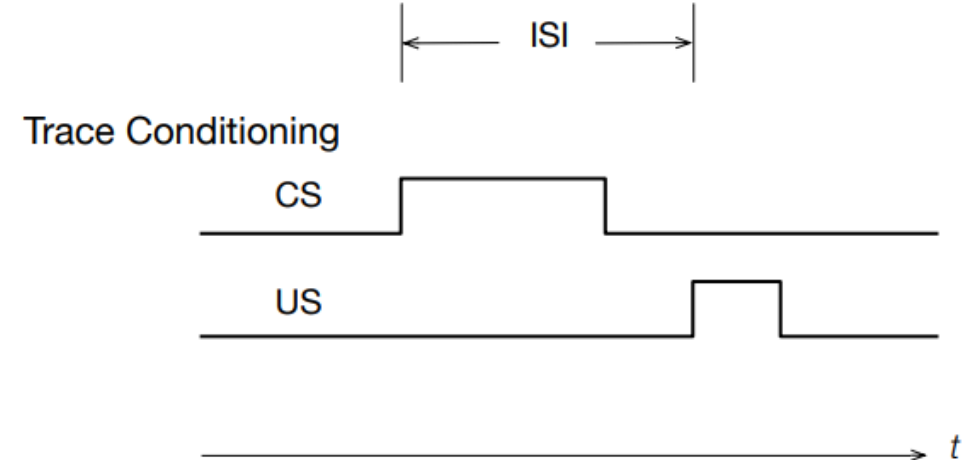
Delay conditioning

- the CS extends throughout the interstimulus interval (ISI), which is the time interval between the CS onset and the US onset (often with the CS and US ending at the same time as shown here).



Trace conditioning

- there is a time interval, called the trace interval, between CS offset and US onset.
- Assumes that a trace of the CS remains when the US arrives. So, learning occurs through the simultaneous presence of the trace and the US
- Learning does not occur across long trace intervals**



Temporal contiguity: the temporal relation between the CS and US is critical for learning

Delay conditioning

- the CS extends throughout the interstimulus interval (ISI), which is the time interval between the CS onset and the US onset (often with the CS and US ending at the same time as shown here).

Trace conditioning

- there is a time interval, called the trace interval, between CS offset and US onset.
- Assumes that a trace of the CS remains when the US arrives. So, learning occurs through the simultaneous presence of the trace and the US
- **Learning does not occur across long trace intervals**

**How would
you test
this?**



Evidence for contiguity learning

Hypothesis: When this CS-US interval is lengthened, a decrement in conditioning is observed. Meaning that it takes more trials for the conditioned response (CR) to appear, and CR strength is often reduced


Experiment

2 groups of rats were exposed to a 6 s tone CS, followed by pellet delivery after:

6s (red data)

18s (purple data)

Anticipatory head entries into the feeding hopper were recorded



Balsam, P. D., Drew, M. R., & Gallistel, C. R. (2010). Time and Associative Learning. *Comparative cognition & behavior reviews*, 5, 1–22.
<https://doi.org/10.3819/ccbr.2010.50001>



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA
CAMPUS DI CESENA

Evidence for contiguity learning

Hypothesis: When this CS-US interval is lengthened, a decrement in conditioning is observed. Meaning that it takes more trials for the conditioned response (CR) to appear, and CR strength is often reduced

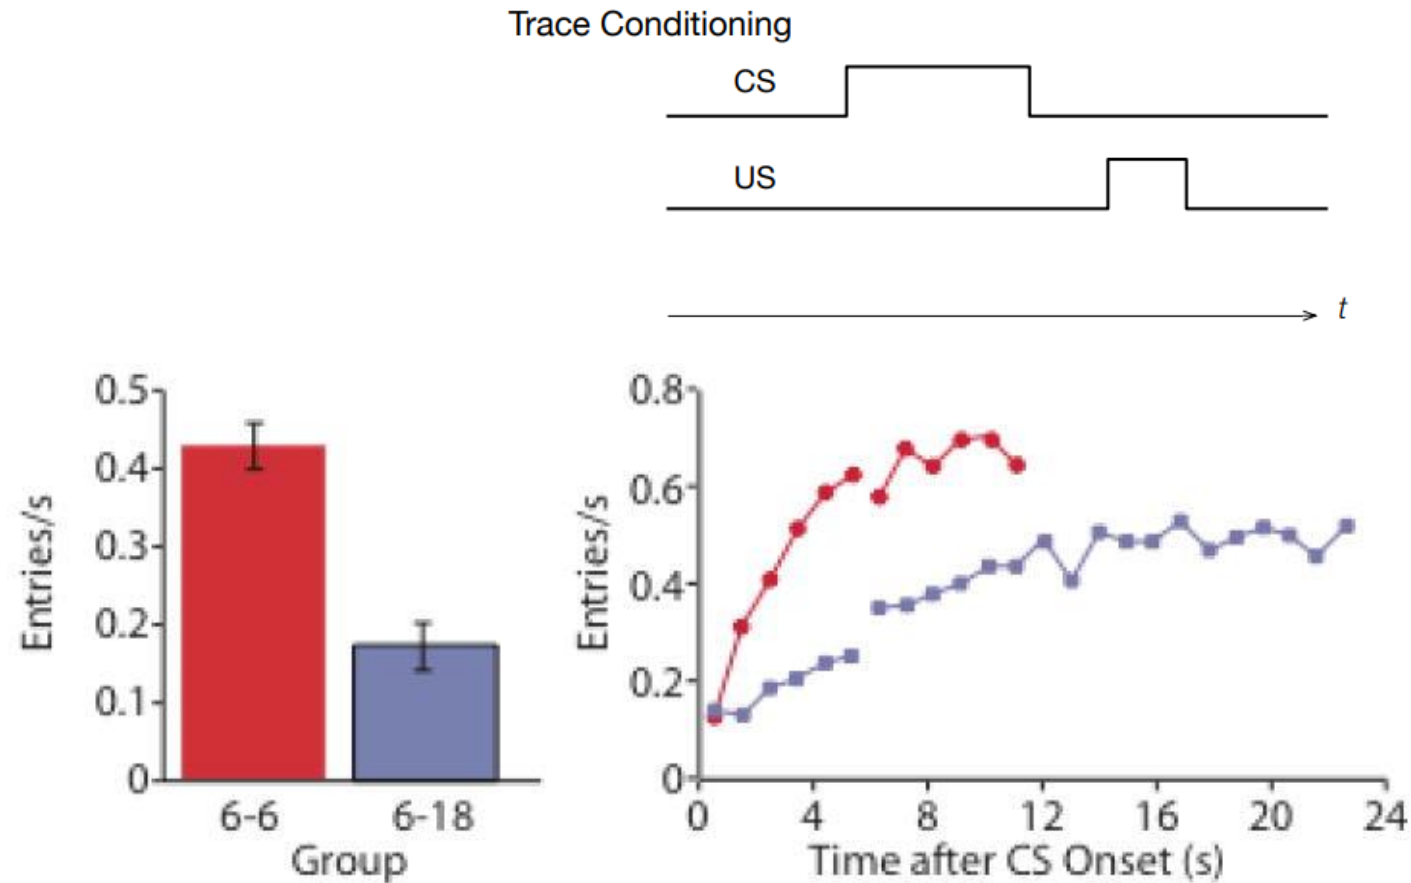
Experiment

2 groups of rats were exposed to a 6 s tone CS, followed by pellet delivery after:

6s (red data)

18s (purple data)

Anticipatory head entries into the feeding hopper were recorded



Balsam, P. D., Drew, M. R., & Gallistel, C. R. (2010). Time and Associative Learning. *Comparative cognition & behavior reviews*, 5, 1–22.
<https://doi.org/10.3819/ccbr.2010.50001>

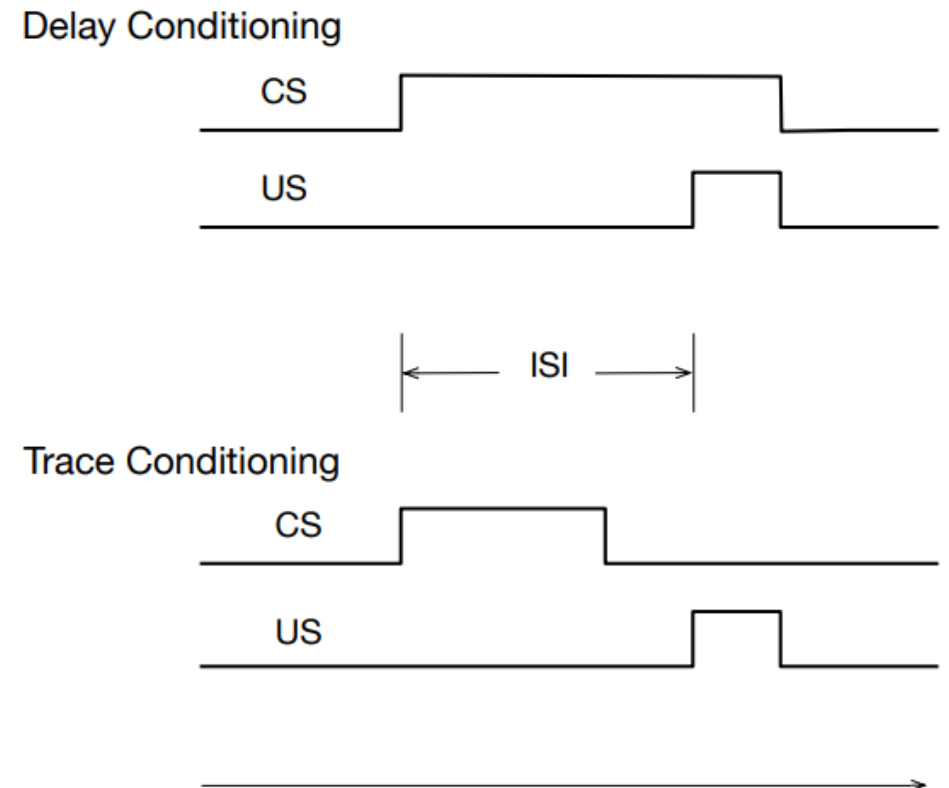


ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA
CAMPUS DI CESENA

Early theories of Pavlovian conditioning took for granted that CS-US contiguity was the critical determinant of learning

Contiguous relations are important because **only closely spaced presentations of the CS and US allow for their mental representations to be simultaneously active and learned about.**

If a long trace interval is introduced, then the CS representation will have decayed by the time the US is presented and thus the CS will not be learned about. (McNally & Westbrook, 2006)




Contiguity learning in the nervous system: the case of Hebbian plasticity

Hebb's Law: **Neurons that fire together, wire together**

All that is needed to condition the gill withdrawal reflex is the simultaneous activation of the syphon and tail neuronal pathways


B Paired pathway




Associative learning is learning about contiguity ...

Contiguity: closeness in time between stimulus/behavior and outcome. Stimuli that are close to one another in time become associated

Is contiguity both necessary and sufficient?





US (dots) contingent on CS



US (dots) not contingent on CS
(but CS-US pairing unchanged)



Challenging the contiguity assumption: The truly random control

In the 1960's and 70's, evidence began to accumulate that posed a challenge to the simple contiguity assumption.


- The truly random control paradigm (Rescorla, 1968)

Showed that **temporal contiguity** between a conditioned stimulus (CS) and an unconditioned stimulus (US) **is not sufficient to lead to learning**



Challenging the contiguity assumption: The truly random control

Experiment (Rescorla, 1968)




The temporal pairing of CS and US is identical in the two groups, but there is no CS-US **contingency** in the second group (the truly random control), because the US occurs as frequently in the absence of the CS as in its presence. The subjects in the Group 1 develop a conditioned response to the CS; the subjects in Group 2 do not.



Challenging the contiguity assumption: The truly random control

Group A


In this experiment, the rat would experience periods of time with a bell, and periods of time without it.



With or without the bell, shock occurs 40% of the time.




Result:
No conditioning to the bell.



Group B


Again, sometimes the bell is sounded, and sometimes not.



As with Group A, the bell is followed by shock 40% of the time. But shock is **less** likely when there is no bell.



Result:
Conditioned fear of the bell.



7.11 The effect of contingency on classical conditioning For both groups, there's only a 40% chance that bells will be followed by shock. However, for Group B, shock is *less likely* when no bell is sounded, and, for this group, the bell becomes a fearful stimulus.

Challenging the contiguity assumption: The truly random control

Contiguity is the same:

The two groups experienced the same number of bell-shock pairings.

Contingency differ:

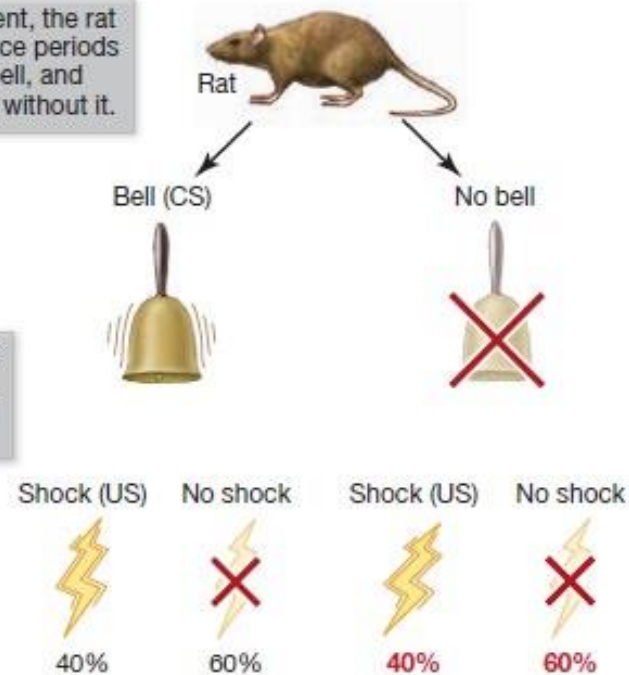
For group B the shock is more likely with the bell than without the bell

Whether the bell is **informative** or not is what matters for conditioning.


Group A

In this experiment, the rat would experience periods of time with a bell, and periods of time without it.

With or without the bell, shock occurs 40% of the time.




Result:
No conditioning to the bell.




Group B

Again, sometimes the bell is sounded, and sometimes not.

As with Group A, the bell is followed by shock 40% of the time. But shock is **less** likely when there is no bell.




Result:
Conditioned fear of the bell.



7.11 The effect of contingency on classical conditioning For both groups, there's only a 40% chance that bells will be followed by shock. However, for Group B, shock is *less likely* when no bell is sounded, and, for this group, the bell becomes a fearful stimulus.

Challenging the contiguity assumption: The truly random control

- The temporal contiguity between the predictor (the CS) and the predicted (the US) is not the key
- but rather the information that the predictor provided about the predicted event



7.11 The effect of contingency on classical conditioning For both groups, there's only a 40% chance that bells will be followed by shock. However, for Group B, shock is *less likely* when no bell is sounded, and, for this group, the bell becomes a fearful stimulus.

Associative learning is more than learning about contiguity. It's also learning about contingency

- Contingency: $p(US/CS) > p(US/no-CS)$
- Learning about **predictive relations** involves learning about more than just the contiguous pairing between the CS and US.
- It involves learning about the **causal relationship** between the two events (Dickinson 1980; Rescorla 1988).



Associative learning is learning about contiguity and contingency...

Contiguity: closeness in time between stimulus/behavior and outcome. Stimuli that are close to one another in time become associated. Not sufficient for learning.

Contingency: causal relationship between stimulus/behavior and outcome. When one stimulus depends on the other, they will become associated --> Predictive value critical





Is it all?



Stage 1

The hiss is reliably followed by the availability of a sexually receptive mate. A CR is thus quickly established.

Hissing noise (CS)

§§§§§

then

Sexually receptive mate (US)



Sexually aroused rat



Stage 2

The procedure continues, but now a light turns on at the same time as the hiss. The light is thus reliably followed by the availability of a mate. This seems like a CS (light) followed by a US (available mate), so it should therefore produce conditioning.

Hissing noise (CS)

§§§§§

+

Light



then

Sexually receptive mate



7.14 The blocking effect procedure

Stage 3

Light



We learn only from surprising/unpredicted events: Blocking

No learning occurs to a stimulus if it is combined with a previously conditioned stimulus during conditioning trials.

The light provided only redundant information. The rat already knew from the hissing noise that the US was about to be presented, and so it wasn't at all surprised when the US did arrive.

$$P(US \mid [\text{compound CS}]) \neq P(US \mid \text{original CS})$$

7.14 The blocking effect procedure

What prevented the conditioning in this sequence? In Stage 2, the light provided no new information, because the hiss told the animal that the US would soon arrive. Conditioning does not occur with uninformative stimuli!

Stage 1

The hiss is reliably followed by the availability of a sexually receptive mate. A CR is thus quickly established.

Hissing noise (CS)

§§§§§

then

Sexually receptive mate (US)



→

Sexually aroused rat



Stage 2

The procedure continues, but now a light turns on at the same time as the hiss. The light is thus reliably followed by the availability of a mate. This seems like a CS (light) followed by a US (available mate), so it should therefore produce conditioning.

Hissing noise (CS)

§§§§§

+

Light



then

Sexually receptive mate



Stage 3

Now we see that conditioning has not occurred: The animal doesn't respond (produces no CR) to the light.

Light



→

Rat shows no sign of sexual arousal

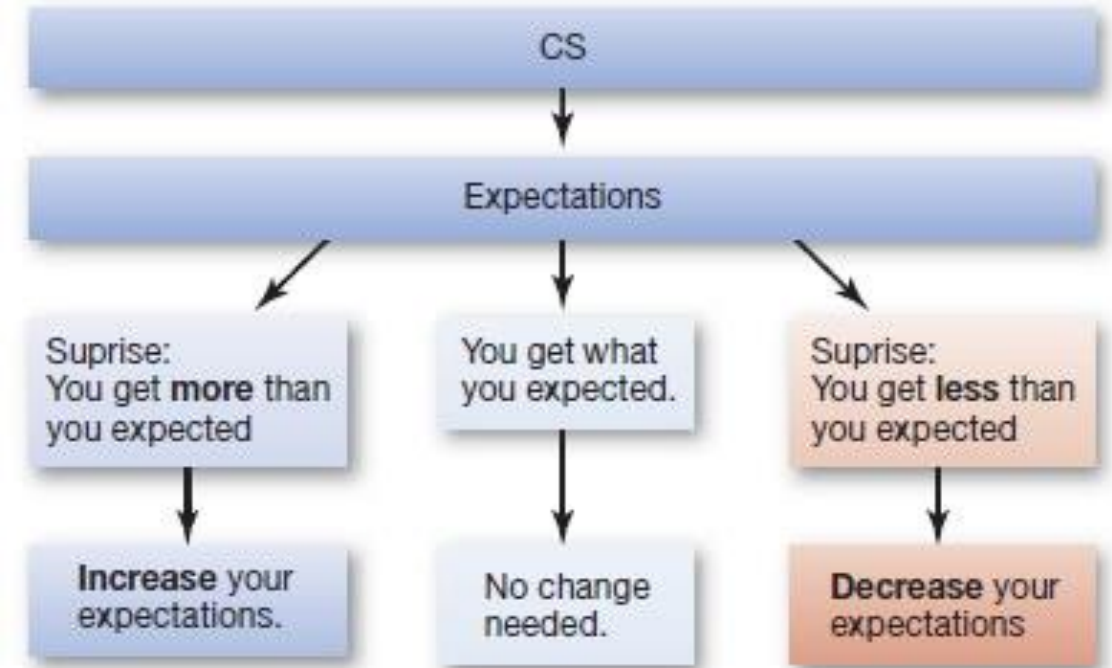


Both contiguity and contingency are not sufficient to learn, you also need surprise



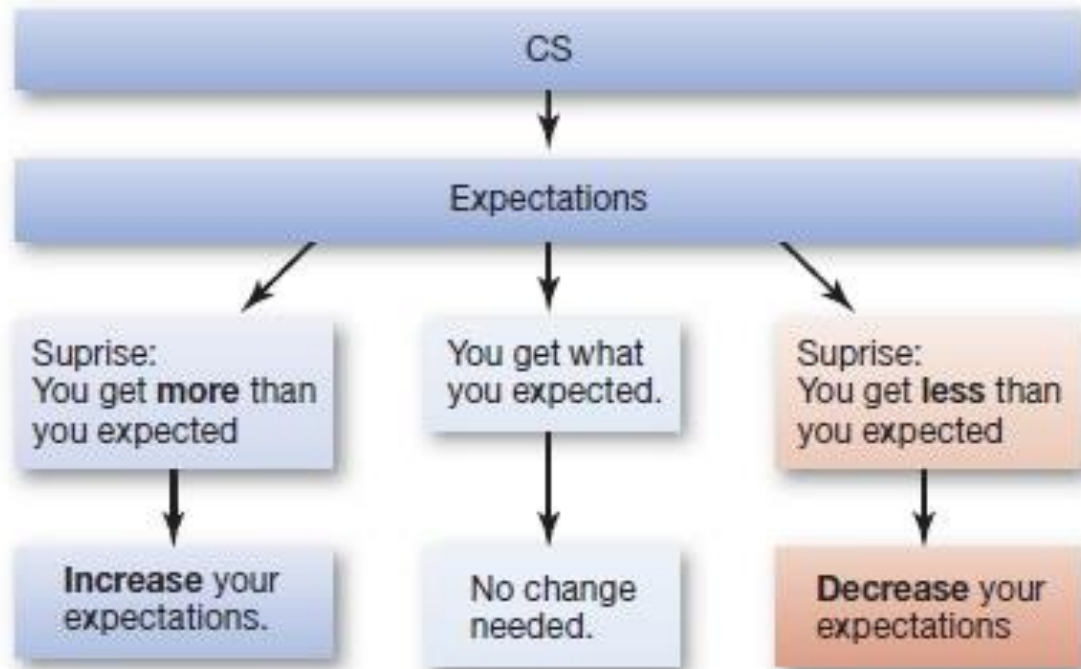
We learn only from surprising/unpredicted events

- Predictive learning depends on what is already known about those events.
- If little is known about the relation between the events, so that the US is not predicted by the CS, then learning occurs.
- If much is known about this relation, so that the US is adequately predicted by the CS, then learning fails.




7.13 The role of expectations and surprise This figure shows the (automatic, unconscious) process through which expectations can be adjusted, trial by trial, in a classical conditioning experiment. The one complication not shown here is that bigger surprises (greater departures from expectations) will trigger larger adjustments; smaller surprises will lead to smaller adjustments.

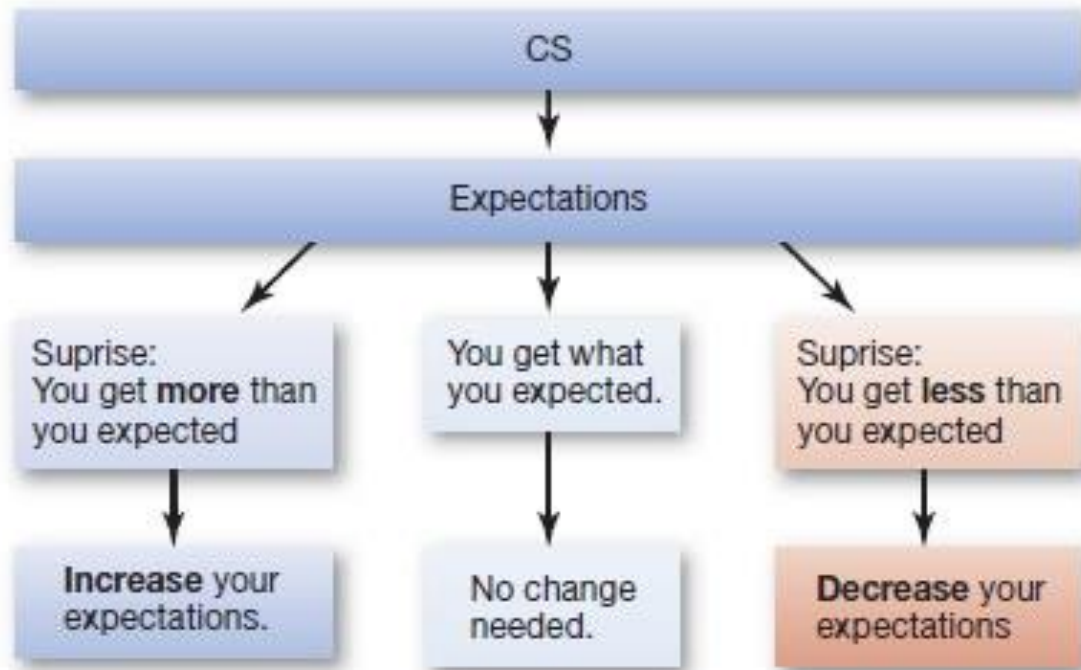
We learn only from surprising/unpredicted events



7.13 The role of expectations and surprise This figure shows the (automatic, unconscious) process through which expectations can be adjusted, trial by trial, in a classical conditioning experiment. The one complication not shown here is that bigger surprises (greater departures from expectations) will trigger larger adjustments; smaller surprises will lead to smaller adjustments.



We learn only from surprising/unpredicted events

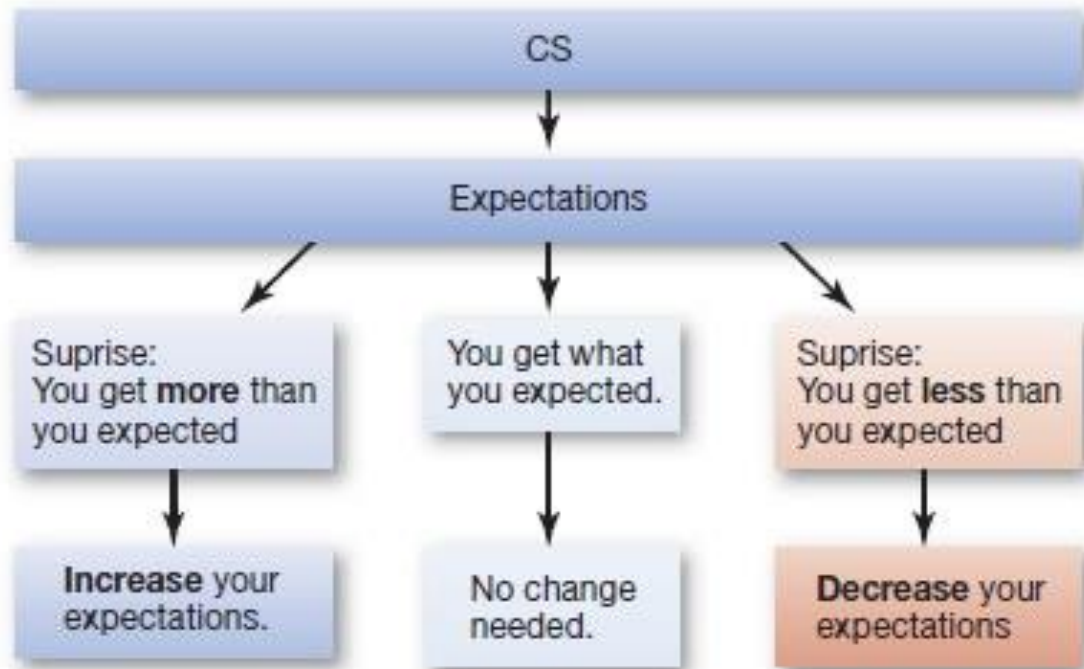


7.13 The role of expectations and surprise This figure shows the (automatic, unconscious) process through which expectations can be adjusted, trial by trial, in a classical conditioning experiment. The one complication not shown here is that bigger surprises (greater departures from expectations) will trigger larger adjustments; smaller surprises will lead to smaller adjustments.

Learning theories have posited that **conditioning is not instructed by a simple sensory representation of the outcome,**

but instead by an **error signal measuring the difference between the outcome actually present and that expected.**
(McNally, 2011)

We learn only from surprising/unpredicted events



7.13 The role of expectations and surprise This figure shows the (automatic, unconscious) process through which expectations can be adjusted, trial by trial, in a classical conditioning experiment. The one complication not shown here is that bigger surprises (greater departures from expectations) will trigger larger adjustments; smaller surprises will lead to smaller adjustments.

Learning is regulated by prediction errors

Prediction error

- is a quantitative discrepancy between the **outcome expected** when the cue was presented, and the **outcome** that was actually **experienced**
- Functions as **teaching signal** to update expectations and reduce following prediction errors



Learning is regulated by prediction errors

- Error signal measuring the difference between the outcome actually present and that expected
- Occurs when the outcome of a conditioning trial is different from that which is predicted by the conditioned stimuli that are present on the trial (i.e., when the US is surprising)

Prediction error is **necessary for associative learning**

- Dictates variations in the effectiveness of the US in supporting learning
- if the difference is large, predictions did not match observations, and there is a need for more learning to update those predictions
- If predictions match observations, there is no prediction error and no learning occurs



Computational models

Mathematical formalization of learning theories



Computational models: the Rescorla-Wagner model

- Change in value is proportional to the difference between predicted and actual outcome --> error driven learning
- R is the value of US
 - R=1 if US is delivered
 - R=0 if US is omitted
- V is the expectancy of US or the expected value of a given CS
- then the error signal is computed as $R - V$
- the expected value V of a given CS is updated based on the sum of the current expected value V and the prediction error at trial t
- Where α is a learning rate parameter, which determines the size of the update step (α assumed between 0 and 1)

Prediction error

$$\delta_t = R_t - V_t$$

$$V_{t+1} = V_t + \alpha \delta_t$$



Excel file




Computational models: the Rescorla-Wagner model

Prototypical acquisition and extinction learning curves for Pavlovian conditioning as predicted by the Rescorla-Wagner model.

The filled circles show the time **evolution of the value V** of a CS over 200 trials.

- In the first 100 trials, a reward R was paired with the CS (acquisition)
- in trials 100-200 no reward was paired (extinction)


Learning is proportional to prediction error δ , which is larger at the start of training when reward is unexpected, and gets smaller, eventually 0, as training progresses and reward is fully predicted.



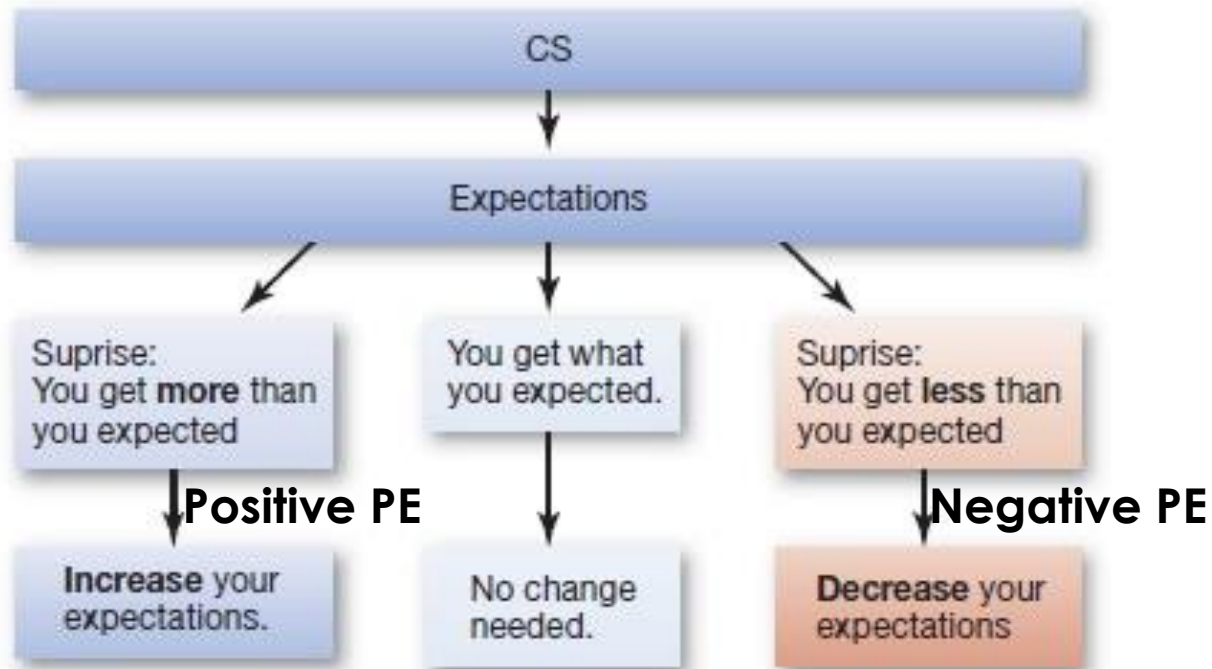
Computational models: the Rescorla-Wagner model

Learning is proportional to prediction error δ


- δ is **positive** (more than expected) and larger at the start of acquisition when reward delivery is unexpected
- δ gets smaller, eventually 0, as acquisition progresses and reward delivery is fully predicted
- δ is **negative** (less than expected) and larger at the start of extinction when reward omission is unexpected
- δ gets smaller, eventually 0, as extinction progresses and reward omission is fully predicted





Rescorla-Wagner (and TD) model rely upon signed prediction errors



7.13 The role of expectations and surprise This figure shows the (automatic, unconscious) process through which expectations can be adjusted, trial by trial, in a classical conditioning experiment. The one complication not shown here is that bigger surprises (greater departures from expectations) will trigger larger adjustments; smaller surprises will lead to smaller adjustments.



Computational models: the Rescorla-Wagner model



Computational models: the Rescorla-Wagner model & blocking

The Rescorla-Wagner model captures blocking by specifying that when multiple stimuli are observed (e.g. light and tone), the animal computes a single net prediction and prediction error, from the sum of all predictions concerning the stimuli present on a given trial.

7.14 The blocking effect procedure

What prevented the conditioning in this sequence? In Stage 2, the light provided no new information, because the hiss told the animal that the US would soon arrive. Conditioning does not occur with uninformative stimuli!

Stage 1

The hiss is reliably followed by the availability of a sexually receptive mate. A CR is thus quickly established.

Hissing noise (CS)

§§§§§

then

Sexually receptive mate (US)



→

Sexually aroused rat



Stage 2

The procedure continues, but now a light turns on at the same time as the hiss. The light is thus reliably followed by the availability of a mate. This seems like a CS (light) followed by a US (available mate), so it should therefore produce conditioning.

Hissing noise (CS)

§§§§§

+

Light



then

Sexually receptive mate



Stage 3

Now we see that conditioning has not occurred: The animal doesn't respond (produces no CR) to the light.

Light



→


Rat shows no sign of sexual arousal



Computational models: the Rescorla-Wagner model


It is a **trial-level** model

- it deals with how associative strengths change from trial to trial without considering any details about what happens within and between trials
- The model does not apply to details about what happens during the time a trial is taking place, or what might happen between trials.
- Within each trial an animal might experience various stimuli whose onsets occur at particular times and that have particular durations. These timing relationships strongly influence learning.




Before conditioning

Food produces salivation, but neither a bell nor a light produces a reaction.




First-order conditioning trials

A pairing of light with food will eventually result in the light triggering salivation.



Second-order conditioning trials



After conditioning




Second-order conditioning

RW model predicts no learning because there is no US in the second-order conditioning trials
BUT
the animal indeed learns

The RW model fails to explain the data


Before conditioning

Food produces salivation, but neither a bell nor a light produces a reaction.




First-order conditioning trials

A pairing of light with food will eventually result in the light triggering salivation.




Second-order conditioning trials

A pairing of bell with light will condition the dog, without ever introducing food.



After conditioning

At the end the bell can trigger salivation.



7.7 Second-order conditioning Before conditioning, neither the bell nor the light trigger salivation. During *first-order* conditioning, the light is paired with meat, and soon presentation of the light can trigger salivation. During *second-order* conditioning, the bell is paired with the light. As a result, presentation of the bell alone (which has never been paired with meat) will elicit salivation.

Computational models: the temporal difference model

From the point of view of the animal, a trial is just a fragment of its continuing experience interacting with its world


- It is a real-time model
- t now labels time steps within a trial instead of complete trials
- **provide a prediction, for each time t in the trial during which a CS is presented, of the total future reward that will be gained in the trial from time t to the end of the trial**
- the prediction error here
 - compares the predicted value at time t to the predicted value at time $t-1$

$$V_{t+1} = V_t + \alpha \delta_t$$

Prediction error

$$\delta_t = R_t + V_t - V_{t-1}$$

Delay Conditioning



Computational models: the temporal difference model

From the point of view of the animal, a trial is just a fragment of its continuing experience interacting with its world

- At the **beginning** of learning, the predicted reward **$V(t)$ is zero for each time t until** the time at which the reward or UCS is delivered (t_{UCS})
- On the **next learning trial**, a comparison between **$V(t_{UCS})$ and $V(t_{UCS} - 1)$** generates a positive prediction error that is used to **increment the V at time $t_{UCS} - 1$** (in proportion to an arbitrary learning rate).
- On **subsequent learning trials**, **$V(t)$ is updated for each time t ranging from t_{UCS} back to t_{CS}** (the earliest time at which the CS is presented).

$$V_{t+1} = V_t + \alpha \delta_t$$

Prediction error

$$\delta_t = R_t + V_t - V_{t-1}$$

Delay Conditioning



Prediction errors in computational models

- **Signed difference between the expected and delivered outcome**
- **Used to update predicted value**
 - The value signal produced by the reward itself transfers back to events that reliably precede reward delivery (e.g. CSs)
 - Thus, **the rewarding value transfers from the reward to the CS that predicts reward**

In this manner, **reinforcement learning algorithms explicitly state that the quantitative value inherent in reward transfers back to the antecedent cue predicting its delivery.** That is, the predictive cue becomes endowed with the scalar value of the reward.






We make predictions all the time

- Predictions are made also in the realm of sensory perception
- Sensory perception is actively constructed
- Sensory perception reflects the brain's prediction about the causes of sensory signals


Are square A and square B of the same or different color?



Edward H. Adelson


We make predictions all the time

Are square A and
square B of the
same or different
color?



Edward H. Adelson

Let's test our
hypothesis



We make predictions all the time

Sensory perception
is
active
and
predictive




A drawback of the predictive brain...



<https://youtu.be/NCtF4aVlxgU>


A drawback of the predictive brain...



Skinner, B. F. (1948). 'Superstition' in the pigeon. *Journal of Experimental Psychology*, 38(2), 168–172. <https://doi.org/10.1037/h0055873>

Abstract

"A pigeon is brought to a stable state of hunger by reducing it to 75 percent of its weight when well fed. It is put into an experimental cage for a few minutes each day. A food hopper attached to the cage may be swung into place so that the pigeon can eat from it. A solenoid and a timing relay hold the hopper in place for five sec. at each reinforcement. If a clock is now arranged to present the food hopper at regular intervals *with no reference whatsoever to the bird's behavior*, operant conditioning usually takes place." The bird tends to learn whatever response it is making when the hopper appears. The response may be extinguished and reconditioned. "The experiment might be said to demonstrate a sort of superstition. The bird behaves as if there were a causal relation between its behavior and the presentation of food, although such a relation is lacking." (PsycINFO Database Record (c) 2016 APA, all rights reserved)



<https://youtu.be/NCtF4aVlxgU>

A drawback of the predictive brain...




A drawback of the predictive brain...

Number of people who drowned by falling into a pool

correlates with

Films Nicolas Cage appeared in




<https://www.tylervigen.com/spurious-correlations>

A drawback of the predictive brain...



Correlation is not causation!



Recommended readings

- Daw, N. D., & Tobler, P.N. (2014). Value Learning through Reinforcement: The Basics of Dopamine and Reinforcement Learning. In *Neuroeconomics* (Chapter 15, pp. 283-298). Academic Press.
- Sutton, R. S., & Barto, A. G. (2018). Reinforcement learning: An introduction. MIT press.
 - Chapter 14
 - Sections 14.1, 14.2

