

Fundamentals of AI and KR - Module 3

5. Approximate inference

Paolo Torroni

Fall 2024

Notice

Downloading and sharing

A copy of these slides can be downloaded from [virtuale](#) and stored for personal use only. Please do not redistribute.

Table of Contents

- About the exam
- Approximate inference

About the exam

Exam formats

Choose one:

- Written exam
- Project

Exam formats

Choose one:

- Written exam
 1. [8 points] one **exercise** in the style of those seen in class or proposed by Russel & Norvig, Koller & Friedman; structured in several points/questions (see Virtuale for past exam exercises)
 2. [3 points] one or more **open questions** about theory, material covered in classes
- Lasts about 1 hour
- How to sit exam: **sign up via AlmaEsami** (be mindful of deadlines)
- 4 dates each year. In 2025:
 - 17 January, 9-11.30
 - 10 February, 9a-11.30
 - 2 July, 9-11.30
 - 11 September, 14.30-17
- Project

Exam formats

Choose one:

- Written exam
- Project
 - Usually two types: vertical case study, or domain-independent investigation of one or more aspects of Bayes nets
 - **Case study:** implement an (interesting) example in a domain of your choice using **pgmpy** or other libraries, explore the model to some extent, for example by studying structural properties, developing queries, making experiments to see what happens if you change parameter values, evidence, etc
 - **Investigation:** choose one or more aspects of Bayes nets you wish to study, make hypotheses, run experiments to validate

Exam formats

Choose one:

- Written exam
- Project
 - Usually two types: vertical case study, or domain-independent investigation of one or more aspects of Bayes nets
 - **Case study**
 - Basic project: existing network (online library, tutorial, textbook, ...), run experiments following examples seen in class
 - Advanced project: original network (own idea, inspired from scientific article, ...), pose interesting questions, develop experiments to answer them, draw motivated conclusions
 - **Investigation**
 - Examples: modeling aspects, theoretical properties, inference methods, learning, software/libraries, ...

Exam formats

Choose one:

- Written exam
- Project
 - Must deliver **code** and 2-page **report** showing aim, model, methods, results (see [guidelines & template](#))
 - Report must follow a **fixed format**; using **L^AT_EX** is encouraged
 - Oral discussion: 10 minutes presentation + 10 minutes Q&A, may include questions on topics seen in class
 - Evaluation based on work done [*4 points*], report [*2 pts*], presentation [*2 pts*] and Q&A [*2 pts*], plus *1 point* for outstanding work (honours)
 - Capacity limited to 7 project discussions each date
 - How to book a date: **sign up via AlmaEsami after uploading project+report** via [link](#) (be mindful of **limited capacity**)
 - There will be dates in January, February, June, July, and September; upon request, also in March and/or May

Approximate inference



Car insurance network




Figure 13.9 A Bayesian network for evaluating car insurance applications.



Car insurance network

Some observations

- Introduction of hidden variables (brown) yields sparse model with reasonable number of parameters
- What causal factors lead where is important question driving modeling choices (topology and hidden variables)
- Another modeling decision is continuous vs discrete ranges
 - Continuous: more precision but exact inference generally impossible
 - Discrete: many values generally yield higher inference cost, unless variable is observed (example: MakeModel)
- Cost of inference with exact enumeration: $\approx 10^8$ operations
- With variable elimination and good ordering: $\approx 10^5$ operations



Inference by stochastic simulation

Basic idea:

- 1 Draw N samples from a sampling distribution S
- 2 Compute an approximate posterior probability $\hat{P}$
- 3 Show this converges to the true probability P



Outline: ^(NO OBSERVATIONS)
WE ASSUME WE KNOW NOTHING

- Sampling from an empty network
- Rejection sampling: reject samples disagreeing with evidence
- Likelihood weighting: use evidence to weight samples
- Markov chain Monte Carlo (MCMC): sample from a stochastic process whose stationary distribution is the true posterior



Performance of approximation algorithms

△ WE PAY EFFICIENCY WITH CORRECTNESS

Absolute approximation: $|P(X|\mathbf{e}) - \hat{P}(X|\mathbf{e})| \leq \epsilon$

Relative approximation: $\frac{|P(X|\mathbf{e}) - \hat{P}(X|\mathbf{e})|}{P(X|\mathbf{e})} \leq \epsilon$

- Relative $\Rightarrow$ absolute since $0 \leq P \leq 1$
- Randomized algorithms may fail with probability at most δ

Theorem (Dagum and Luby, 1993)

Both absolute and relative approximation for either deterministic or randomized algorithms are NP-hard for any $\epsilon, \delta < 0.5$




Sampling from an empty network

```
function PRIOR-SAMPLE(bn) returns an event sampled from bn  
  inputs: bn, a belief network specifying joint distribution  $\mathbf{P}(X_1, \dots, X_n)$   
   $\mathbf{x} \leftarrow$  an event with  $n$  elements  
  for  $i = 1$  to  $n$  do  
     $x_i \leftarrow$  a random sample from  $\mathbf{P}(X_i | \text{parents}(X_i))$   
    given the values of  $\text{Parents}(X_i)$  in  $\mathbf{x}$   
  return  $\mathbf{x}$ 
```

Idea: sample each variable in turn, in topological order




Example






Example






Example






Example






Example





Example





Sampling from an empty network

Probability that PRIORSAMPLE generates a particular event

$$S_{PS}(x_1 \dots x_n) = \prod_{i=1}^n P(x_i | \text{parents}(X_i)) = P(x_1 \dots x_n)$$

i.e., the true prior probability

$$\text{E.g., } S_{PS}(t, f, t, t) = 0.5 \times 0.9 \times 0.8 \times 0.9 = 0.324 = P(t, f, t, t)$$

Let $N_{PS}(x_1 \dots x_n)$ be the number of samples generated for event $x_1, \dots, x_n$

Then we have

$$\begin{aligned} \lim_{N \rightarrow \infty} \hat{P}(x_1, \dots, x_n) &= \lim_{N \rightarrow \infty} N_{PS}(x_1, \dots, x_n) / N \\ &= S_{PS}(x_1, \dots, x_n) \\ &= P(x_1 \dots x_n) \end{aligned}$$

That is, estimates derived from PRIORSAMPLE are **consistent**

Shorthand: $\hat{P}(x_1, \dots, x_n) \approx P(x_1 \dots x_n)$



Rejection sampling VERY INEFFICIENT

Is like the previous method but you change your range of focus only to the samples that present the evidence and reject the others

$\hat{P}(X|e)$ estimated from samples agreeing with e

function REJECTION-SAMPLING(X, e, bn, N) **returns** an estimate of $P(X|e)$

local variables: N , a vector of counts over X , initially zero

for $j = 1$ **to** N **do**

$x \leftarrow \text{PRIOR-SAMPLE}(bn)$

if x is consistent with e **then**

$N[x] \leftarrow N[x] + 1$ where x is the value of X in x

return NORMALIZE(N)

E.g., estimate $P(\text{Rain} | \text{Sprinkler} = \text{true})$ using 100 samples

27 samples have $\text{Sprinkler} = \text{true}$ (not $\text{sprinkler} = \text{false}$)

Of these, 8 have $\text{Rain} = \text{true}$ and 19 have $\text{Rain} = \text{false}$.

$\hat{P}(\text{Rain} | \text{Sprinkler} = \text{true}) = \text{NORMALIZE}(\langle 8, 19 \rangle) = \langle 0.296, 0.704 \rangle$

Similar to a basic real-world empirical estimation procedure



Analysis of rejection sampling

$$\begin{aligned}
 \hat{\mathbf{P}}(X|\mathbf{e}) &= \alpha \mathbf{N}_{PS}(X, \mathbf{e}) && \text{(algorithm defn.)} \\
 &= \mathbf{N}_{PS}(X, \mathbf{e}) / N_{PS}(\mathbf{e}) && \text{(normalized by } N_{PS}(\mathbf{e})\text{)} \\
 &\approx \mathbf{P}(X, \mathbf{e}) / P(\mathbf{e}) && \text{(property of PRIORSAMPLE)} \\
 &= \mathbf{P}(X|\mathbf{e}) && \text{(defn. of conditional probability)}
 \end{aligned}$$

Hence rejection sampling returns consistent posterior estimates

Problem: hopelessly expensive if $P(\mathbf{e})$ is small

$P(\mathbf{e})$ drops off exponentially with number of evidence variables!

TOO GREEDY!!



Likelihood weighting

Idea: fix evidence variables, sample only nonevidence variables and weight each sample by the likelihood it accords the evidence

```

function LIKELIHOOD-WEIGHTING( $X, e, bn, N$ ) returns an estimate of  $P(X|e)$ 
  local variables:  $W$ , a vector of weighted counts over  $X$ , initially zero
  for  $j = 1$  to  $N$  do
     $x, w \leftarrow$  WEIGHTED-SAMPLE( $bn, e$ )
     $W[x] \leftarrow W[x] + w$  where  $x$  is the value of  $X$  in  $x$ 
  return NORMALIZE( $W[X]$ )
  
```


```

function WEIGHTED-SAMPLE( $bn, e$ ) returns an event and a weight
   $x \leftarrow$  an event with  $n$  elements initialized from  $e$ 
   $w \leftarrow 1$ 
  for  $i = 1$  to  $n$  do
    if  $X_i$  is an evidence variable with value  $x_i$  in  $e$  then
       $w \leftarrow w \times P(X_i = x_i | \text{parents}(X_i))$ 
    else  $x_i \leftarrow$  a random sample from  $P(X_i | \text{parents}(X_i))$ 
  return  $x, w$ 
  
```




Likelihood weighting example $P(x|S, W_G)$?

We generate random values in $[0,1]$ and keeping previous numbers: 0.4, 0.7, 0.1, 0.9



HERE WE USE THE RANDOM VALUES BECAUSE WE HAVE TO ASSIGN VALUES



HERE WE DON'T LOOK AT RANDOM NUMBERS SINCE WE WANT IT TRUE

$W_S \leftarrow$

C	$P(S C)$
T	.10
F	.50

Sprinkler

Rain

C	$P(R C)$
T	.80
F	.20

WetGrass

$W_{W_G} \leftarrow$

S	R	$P(W S, R)$
T	T	.99
T	F	.90
F	T	.90
F	F	.01

INITIALLY


$$w = \overline{1.0} \cdot W_S \cdot W_{W_G}$$

⚠ NOT " W_R " SINCE RAIN HAS BEEN SAMPLED

THEN WE DON'T DIVIDE BY THE NUMBER OF SAMPLES BUT WE DIVIDE BY THE SUM OF ALL THE WEIGHTS
 $\rightarrow$ they keep increasing and don't sum up to 1




Likelihood weighting example



$w = 1.0$




Likelihood weighting example



$$w = 1.0$$



Likelihood weighting example




$$w = 1.0$$

$$w = 1.0 \times 0.1$$



Likelihood weighting example




$$w = 1.0$$

$$w = 1.0 \times 0.1$$



Likelihood weighting example




$$w = 1.0$$

$$w = 1.0 \times 0.1$$



Likelihood weighting example



$$w = 1.0$$

$$w = 1.0 \times 0.1$$

$$w = 1.0 \times 0.1 \times 0.99 = 0.099$$




Analysis of likelihood weighting

Sampling probability for WEIGHTEDSAMPLE is

$$S_{WS}(\mathbf{z}, \mathbf{e}) = \prod_{i=1}^I P(z_i | \text{parents}(Z_i))$$

Note: pays attention to evidence in ancestors only

⇒ somewhere “in between” prior and posterior distribution



Weight for a given sample $\mathbf{z}, \mathbf{e}$ is $w(\mathbf{z}, \mathbf{e}) = \prod_{i=1}^m P(e_i | \text{parents}(E_i))$

Weighted sampling probability is $S_{WS}(\mathbf{z}, \mathbf{e})w(\mathbf{z}, \mathbf{e})$


$$\begin{aligned}
 &= \prod_{i=1}^I P(z_i | \text{parents}(Z_i)) \prod_{i=1}^m P(e_i | \text{parents}(E_i)) \\
 &= P(\mathbf{z}, \mathbf{e}) \text{ (by standard global semantics of network)}
 \end{aligned}$$

Hence likelihood weighting returns consistent estimates but performance still degrades with many evidence variables because a few samples have nearly all the total weight



Markov chain Monte Carlo example

With *Sprinkler* = true, *WetGrass* = true, there are four states:



Wander about for a while, average what you see



Markov chain Monte Carlo example


Estimate $\mathbf{P}(\text{Rain} | \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true})$

Sample *Cloudy* or *Rain* given its Markov blanket, repeat.

Count number of times *Rain* is true and false in the samples.

E.g., visit 100 states, 31 have *Rain* = true, 69 have *Rain* = false

$$\hat{\mathbf{P}}(\text{Rain} | \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true}) \\ = \text{NORMALIZE}(\langle 31, 69 \rangle) = \langle 0.31, 0.69 \rangle$$





The Markov chain

“State” of network = current assignment to all variables.

Generate next state by sampling one variable given Markov blanket

Sample each variable in turn, keeping evidence fixed

Theorem

*Theorem: chain approaches **stationary distribution**: long-run fraction of time spent in each state is exactly proportional to its posterior probability*



Markov chain Monte Carlo

```

function MCMC-ASK( $X, e, bn, N$ ) returns an estimate of  $P(X|e)$ 
  local variables:  $N$ , a vector of counts for each value of  $X$ , initially zero
                     $Z$ , the nonevidence variables in  $bn$ 
                     $x$ , the current state of the network, initially copied from  $e$ 

  initialize  $x$  with random values for the variables in  $Z$ 
  for  $j = 1$  to  $N$  do
    for each  $Z_i$  in  $Z$  do
      set the value of  $Z_i$  in  $x$  by sampling from  $P(Z_i|mb(Z_i))$ 
       $N[x] \leftarrow N[x] + 1$  where  $x$  is the value of  $X$  in  $x$ 
  return NORMALIZE( $N$ )
  
```


Can also choose a variable to sample at random each time



Markov blanket sampling

Markov blanket of *Cloudy* is ...

Markov blanket of *Rain* is ...



Probability given the Markov blanket is calculated as follows:

$$P(x'_i | mb(X_i)) = P(x'_i | parents(X_i)) \prod_{Z_j \in Children(X_i)} P(z_j | parents(Z_j))$$

Main computational problems:


- 1 Difficult to tell if convergence has been achieved
- 2 Can be wasteful if Markov blanket is large:
 $P(X_i | mb(X_i))$ won't change much (law of large numbers)



Markov blanket sampling

Markov blanket of *Cloudy* is
 $\{\textit{Sprinkler}, \textit{Rain}\}$

Markov blanket of *Rain* is
 $\{\textit{Cloudy}, \textit{Sprinkler}, \textit{WetGrass}\}$



Probability given the Markov blanket is calculated as follows:

$$P(x'_i | mb(X_i)) = P(x'_i | parents(X_i)) \prod_{Z_j \in Children(X_i)} P(z_j | parents(Z_j))$$

Main computational problems:

- 1 Difficult to tell if convergence has been achieved
- 2 Can be wasteful if Markov blanket is large:
 $P(X_i | mb(X_i))$ won't change much (law of large numbers)



MCMC analysis: Outline

- Transition probability $q(\mathbf{x} \rightarrow \mathbf{x}')$
- Occupancy probability $\pi_t(\mathbf{x})$ at time t
- Equilibrium condition on π_t defines stationary distribution $\pi(\mathbf{x})$
Note: stationary distribution depends on choice of $q(\mathbf{x} \rightarrow \mathbf{x}')$
- Pairwise **detailed balance** on states guarantees equilibrium
- **Gibbs sampling** transition probability: sample each variable given current values of all others
 - $\Rightarrow$ detailed balance with the true posterior
- For Bayesian networks, Gibbs sampling reduces to sampling conditioned on each variable's Markov blanket



Stationary distribution

$\pi_t(\mathbf{x})$ = probability in state $\mathbf{x}$ at time t

$\pi_{t+1}(\mathbf{x}')$ = probability in state $\mathbf{x}'$ at time $t + 1$

π_{t+1} in terms of π_t and $q(\mathbf{x} \rightarrow \mathbf{x}')$

$$\pi_{t+1}(\mathbf{x}') = \sum_{\mathbf{x}} \pi_t(\mathbf{x}) q(\mathbf{x} \rightarrow \mathbf{x}')$$

Stationary distribution: $\pi_t = \pi_{t+1} = \pi$

$$\pi(\mathbf{x}') = \sum_{\mathbf{x}} \pi(\mathbf{x}) q(\mathbf{x} \rightarrow \mathbf{x}') \quad \text{for all } \mathbf{x}'$$

If π exists, it is unique (specific to $q(\mathbf{x} \rightarrow \mathbf{x}')$)

In equilibrium, expected “outflow” = expected “inflow”



Detailed balance

“Outflow” = “inflow” for each pair of states:

$$\pi(\mathbf{x})q(\mathbf{x} \rightarrow \mathbf{x}') = \pi(\mathbf{x}')q(\mathbf{x}' \rightarrow \mathbf{x}) \quad \text{for all } \mathbf{x}, \mathbf{x}'$$

Detailed balance $\Rightarrow$ stationarity:

$$\begin{aligned} \sum_{\mathbf{x}} \pi(\mathbf{x})q(\mathbf{x} \rightarrow \mathbf{x}') &= \sum_{\mathbf{x}} \pi(\mathbf{x}')q(\mathbf{x}' \rightarrow \mathbf{x}) \\ &= \pi(\mathbf{x}') \sum_{\mathbf{x}} q(\mathbf{x}' \rightarrow \mathbf{x}) \\ &= \pi(\mathbf{x}') \end{aligned}$$

MCMC algorithms typically constructed by designing a transition probability q that is in detailed balance with desired π



Gibbs sampling

Sample each variable in turn, given all other variables

Sampling X_i , let $\bar{\mathbf{x}}_i$ be all other nonevidence variables

Current values are x_i and $\bar{\mathbf{x}}_i$; $\mathbf{e}$ is fixed

Transition probability is given by

$$q(\mathbf{x} \rightarrow \mathbf{x}') = q(x_i, \bar{\mathbf{x}}_i \rightarrow x'_i, \bar{\mathbf{x}}_i) = P(x'_i | \bar{\mathbf{x}}_i, \mathbf{e})$$

This gives detailed balance with true posterior $P(\mathbf{x} | \mathbf{e})$:

$$\begin{aligned}\pi(\mathbf{x})q(\mathbf{x} \rightarrow \mathbf{x}') &= P(\mathbf{x} | \mathbf{e})P(x'_i | \bar{\mathbf{x}}_i, \mathbf{e}) = P(x_i, \bar{\mathbf{x}}_i | \mathbf{e})P(x'_i | \bar{\mathbf{x}}_i, \mathbf{e}) \\ &= P(x_i | \bar{\mathbf{x}}_i, \mathbf{e})P(\bar{\mathbf{x}}_i | \mathbf{e})P(x'_i | \bar{\mathbf{x}}_i, \mathbf{e}) \quad (\text{chain rule}) \\ &= P(x_i | \bar{\mathbf{x}}_i, \mathbf{e})P(x'_i, \bar{\mathbf{x}}_i | \mathbf{e}) \quad (\text{chain rule backwards}) \\ &= q(\mathbf{x}' \rightarrow \mathbf{x})\pi(\mathbf{x}') = \pi(\mathbf{x}')q(\mathbf{x}' \rightarrow \mathbf{x})\end{aligned}$$



Compiling approximate inference

- Sampling operates on Bayes net represented as data structure
- Naive implementation would yield huge amounts of repeated operations (example: to find a node's parents)
- Repetition completely unnecessary
- For example, in the alarm network, an operation required by MCMC-ASK is sampling from $P(Z_i | mb(Z_i))$
- Sampling from $P(\text{Earthquake} | mb(\text{Earthquake}))$ requires
 - looking up children and parents of **Earthquake** in the network structure
 - looking up their current values
 - using these values to look up into CPT
 - computing products to obtain distribution
 - sample from distribution
- Solution: *compile* the network into model-specific inference code



Compiling approximate inference

For example:

```
 $r \leftarrow$  a uniform random sample from  $[0,1]$   
if Alarm = true  
  then if Burglary = true  
    then return  $[r < 0.0020212]$   
    else return  $[r < 0.36755]$   
  else if Burglary = true  
    then return  $[r < 0.0016672]$   
    else return  $[r < 0.0014222]$ 
```

Code not pretty but 2/3 orders of magnitude faster than MCMC-Ask



Summary

Exact inference by variable elimination

- only for discrete random variables and very special cases of continuous variables with canonical distributions
- polytime on polytrees, NP-hard on general graphs
- space = time, very sensitive to topology

Approximate inference by LW, MCMC:

- Can handle arbitrary combinations of discrete and continuous variables
- LW does poorly when there is lots of (downstream) evidence
- LW, MCMC generally insensitive to topology
- Convergence can be very slow with probabilities close to 1 or 0
- Compiling could be crucial

Questions?