

CHAPTER 11



The Kantian Perspective

Fairness and Justice

Imagine a person who reasons as follows: I should keep my money rather than pay it out in taxes, because if I keep it, I'll be able to afford a wonderful vacation for myself and my family. And no one is actually going to suffer if I pocket the money, since it's only a few thousand dollars that we're talking about. There's no way that money could bring as much happiness in the government's hands as it could in mine.

Suppose he is right about that. He spends the money on his vacation. He and his family have a terrific time. He is never caught.

Still, he has done something wrong. So has the person who cheats on her exams and gets away with it. So has the person who gleefully speeds down the emergency lane and escapes the traffic jam that the rest of us are stuck in. So has the person whose campaign of dirty tricks has gotten him securely into office.

Despite any good results that may come from their actions, these people did wrong—or so we think. And the explanation of their immorality is simple. What they did was unfair. They took advantage of the system. They broke the rules that work to everyone's benefit. They violated the rights of others. No matter how much personal gain such actions bring, they are still wrong, because they are unfair and unjust.

Immanuel Kant (1724–1804) thought this way, and was very likely the most brilliant philosopher ever to have done so. He remains perhaps the most important voice of opposition to utilitarianism, and to its claim that the ultimate point of morality is to improve well-being rather than do justice.

Consistency and Fairness

There is a natural way to understand what is wrong with the actions in the examples just given. In each case, people are making exceptions of themselves. Their success depends on violating rules that most other people are following. This is a kind of inconsistency—of playing by one set of rules while insisting that others obey a different set.

People are inconsistent to the extent that they treat similar cases differently. Tax cheats or dirty politicians are in the same boat as the rest of us. There's nothing special about them, or their situation, that exempts them from the rules that everyone must follow. That you can get away with making an exception of yourself doesn't mean that it is right to do so.

Our deep opposition to unfairness, and the resulting importance we attach to consistency, are revealed in two very popular tests of morality. Each takes the form of a question:

1. What if everyone did that?
2. How would you like it if I did that to you?

When we ask such questions—in the face of a bully, a liar, or a double-crosser—we are trying to get the person to see that he is acting unfairly, making an exception of himself, living by a set of rules that work only because others are not doing what he is doing. These basic moral challenges are designed to point out the inconsistency, and so the immorality, of that person's behavior.

Consider the first question: What if everyone did that? This question is really shorthand for the following test: *if disastrous results would occur if everyone did X, then X is immoral*. If everyone used the emergency lanes in traffic jams, then ambulances and fire trucks would often fail to provide needed help, leaving many to die. If everyone cheated on their taxes, society would crumble. If every candidate resorted to dirty tricks, then the entire political system would become corrupted. The test works easily and well for these cases.

But the test fails for other cases, and so it cannot serve as a reliable way to learn the morality of actions. Consider a common argument against homosexual sex: if everyone did that, disaster would soon follow, for the human race would quickly die out. Even if this were true, that wouldn't show that homosexual sex is immoral. Why not? Well, consider those who have decided to remain celibate—perhaps they are priests, or committed lifelong

bachelors who believe that one shouldn't have sex without being married. What if everyone did *that*—in other words, refrained from having sex? The same results would follow. But that doesn't show that celibacy is immoral.

The real problem for this test, apart from the fact that it sometimes delivers mistaken verdicts, is that it makes the morality of an action depend on how it is described. Suppose the sexual relations of a gay couple were described as their having consensual, enjoyable sex. In that case, their actions would pass the test. But that undermines the test, because it shows that the test yields contradictory results. The very same action is said to be both morally wrong and morally acceptable, depending only on how it is described. Without any independent guidance on how to select one description over another, this test cannot do the job it was supposed to do—namely, identify which acts are immoral.

What about the other test, the one that asks: How would you like it if I did that to you? This is a direct application of the golden rule, which tells you to treat others as you would like to be treated. The golden rule is the classic test of morality. Clearly, it is meant to be a test of consistency. If you wouldn't want to be slandered or exploited, then don't do such things to others. If you do them anyway, you are acting inconsistently, hence unfairly, and therefore immorally.

Getting people to imagine what it would be like to switch places with their intended victims is often a very effective way to convey a moral message. That is why films and literature are often such powerful tools of moral education. But imaginatively filling someone else's shoes, and asking yourself whether you'd accept being treated in a certain way, is actually an unreliable test of morality. The golden rule cannot be correct.

Kant himself identified the basic reason for this. The golden rule makes morality depend on a person's desires. Most of us don't like to be hit. And so the golden rule forbids us from hitting others. Good. But what about masochists who enjoy being hit? The golden rule allows them to go around hitting others. That's bad. The morality of hitting people shouldn't depend on whether you like to take a beating every now and then.

Consider a related problem, that of the fanatic. Fanatics are principled people. It's just that their principles are ones that we find frightening and revolting. Some fanatics are so wedded to their cause, so strong-willed and self-disciplined, that they would accept the suffering that they want to impose on their victims, were the role of victim and persecutor reversed. True, few Nazis, for instance, would really accept a march to the gas chamber were they to discover their Jewish ancestry. Most Nazis, like most

fanatics generally, are opportunists of bad faith, ones with very limited empathy and only a feeble ability to imagine themselves in someone else's place. If roles really were reversed, they'd much more likely beg for mercy and abandon their genocidal principles. But some would not. There are true believers out there who are willing to suffer any harm in the name of their chosen cause. The golden rule licenses their extremism because it makes the morality of an action depend entirely on what you want and what you are willing to put up with.

The golden rule also fails to give us guidance on **self-regarding actions** (i.e., those that concern only oneself). That's not a problem for most people these days, since it's now unusual to think that we owe moral duties to ourselves. But in Kant's time, self-regarding duties were widely endorsed, and many people still think, for instance, that there is something immoral about suicide or about letting one's talents go to waste, even if no one else is harmed in the process.

Because the golden rule sometimes gives the wrong answer to moral questions, it cannot be the ultimate test of morality. Something else must explain why it works, when it does. Kant thought he had the answer.

The Principle of Universalizability

Kant, like most of us, felt the appeal of the two tests just discussed. He agreed that common sense is deeply committed to the importance of fairness and consistency, something that these two tests were trying, unsuccessfully, to capture. His aim was to identify the ultimate principle of morality, one that would explain the attraction of the two tests while correcting for their shortcomings.

He thought he had found it in the following standard, the **principle of universalizability**:

An act is morally acceptable if, and only if, its maxim is universalizable.

To understand what this means, we need to understand two things: what a **maxim** is, and what it is for a maxim to be **universalizable**.

A **maxim** is simply the principle of action you give yourself when you are about to do something. For instance, if you send a regular check to Oxfam, your maxim might be: contribute \$50 per month to Oxfam to help reduce hunger. A maxim has two parts. It states what you are about to do, and why you are about to do it. You dictate your own maxims. These are the rules you live by.

Kant thought that every action has a maxim. Of course we don't always formulate these maxims clearly to ourselves prior to acting, but at some level, whenever we act, we intend to do something, and we have a reason for doing it. A maxim is nothing but a record of that intention and its underlying reason. Maxims are what we cite when we try to explain to others why we act as we do.

If we lack a maxim, then we aren't really acting at all. We could be moving our bodies, as we do when we sneeze or roll across the bed in our sleep. But the absence of a maxim in these cases shows that these are mere bodily movements, rather than genuine actions.

Kant thought that an action's rightness depends on its maxim. And this leads directly to a very important point. For Kant, the morality of our actions has nothing to do with results. It has everything to do with our intentions and reasons for action, those that are contained in the principles we live by. This is a clear break with consequentialism.

Indeed, we can imagine two people doing the same thing, but for different reasons. That means that they will have different maxims. And even if their actions bring about identical results, one of the actions may be right and the other wrong, since only one of the maxims may be morally acceptable. This is something that act consequentialists cannot accept.

It might be, for instance, that I keep my promises to you because I think it's right to do so. But I might instead keep my promises because I want you to like me so much that you leave your fortune to me in your will. Assume that these different reasons don't change the results of keeping my promises. Then the utilitarian thinks that each case of promise keeping is equally good. But since my maxim is different in these cases, Kant thinks that the morality of these actions might be different. It all depends, as we'll shortly see, on whether their maxims are universalizable.

Many people agree with Kant's view that the morality of our actions depends not on their results, but on our maxims. This supports our thought that those who set out to do evil are acting immorally, even if, through sheer chance, they manage to do good. It also justifies the claim that people who live by noble principles are acting morally, even when some unforeseeable accident intervenes, and their action brings only bad results.

Kant had a deep reason (discussed in detail in the next chapter) for making the morality of an action depend on its maxim rather than its results: it is crucial that the morality of our actions depends entirely on what is within our control. We *can* control which maxims will govern our actions.

We decide for ourselves what we intend to do and what reasons we endorse for doing it. Even in cases where my options are severely limited, as when a thug has a gun to my head, it is up to me to decide which choice to make.

By contrast, the results of our actions are often out of our hands. We can't always control them. And it is unfair to assign credit or blame for things we can't control. That is why we have an insanity defense. That is why we don't prosecute animals for the damage they sometimes cause. That is why we don't condemn infants for any harm they do.

So the morality of actions depends on their maxims. But how, precisely? Not every maxim is going to be a good one. We need a way to sort out the good maxims from the bad. That's where universalizability comes in.

How can we tell whether a maxim is universalizable? Here is a three-part test:

1. Formulate your maxim clearly—state what you intend to do, and why you intend to do it.
2. Imagine a world in which everyone supports and acts on your maxim.
3. Then ask: Can the goal of my action be achieved in such a world?

If the answer to this last question is *yes*, then the maxim is universalizable, and the action is morally acceptable. If the answer is *no*, then the maxim is not universalizable, and the action it calls for is immoral.

This should strike a familiar note. The test of a maxim's universalizability clearly echoes the rule consequentialist's test for optimific social rules (see the previous chapter), and the *what if everyone did that?* test discussed above. Indeed, Kant has us ask a version of that question in the second step of this three-part test. But unlike these other tests, Kant doesn't ask about whether people would be much better off in the imagined world, or about whether disaster would strike there. Instead, he asks about whether we could achieve our goals in that world. But what is so important about that?

The importance, for Kant, is that this three-part test serves as the real way to determine whether we are being consistent and fair. If our maxim is universalizable, then we are pursuing actions for reasons that everyone could stand behind. We are not making exceptions of ourselves. Our goals are ones that everyone *could* support, even if, in the real world, some are dead set against them. We are asking whether our aims could be achieved if everyone shared them. If they can be, this shows that we are living by fair rules. Were we making an exception of ourselves, our maxims wouldn't be universalizable.

Consider the tax cheat again. The only reason he can get what he is aiming for (a lovely vacation) is because enough others are not adopting his maxim. The same goes for the careless driver who speeds down the emergency lane.

Kant sought to make this point with an example of his own. (I am embellishing a bit, but the essence of the example is Kant's.) Suppose that I am a compulsive gambler who is constantly in debt. One night I go to the tables to recoup my losses, only to dig myself further in the hole. The casino boss is having his men drop by tonight to collect. I can either pay them or have my kneecaps broken. I know which one I'd prefer.

The problem is, I don't have the money, the bank won't lend me any more, and I don't have anyone to turn to but you. Since you are aware of my reputation, I know that the only way to get the money is by lying to you. So I beg and plead and promise you, by all I hold dear, that I will repay you—even though I have no intention of doing so. I have just made what Kant calls a *lying promise*.

It seems clear that what I am about to do is immoral. And that is true even if, through a minor miracle, I then feel so much guilt that I repent of my ways, transform myself, and make the lie turn out for the best in the end. The morality of the action doesn't depend on its results, but on its maxim. And my maxim here is not universalizable. So my action is immoral, as Kant says, and as we believe.

Here's why. Suppose my maxim is: lie to a friend, in order to escape from being hurt. And suppose everyone acts on this maxim. They lie whenever they think that it is necessary to avoid some personal harm. In that situation, no one would trust the promises of others. And without that trust, people could not achieve the goals they are aiming for with their promises. In a world where no one believed the promises of others, I'd never be able to get money from you with my promise. And so the purpose of my promise would be defeated. And so my maxim is not universalizable. I am making an exception of myself, and am treating you unfairly. My action is therefore immoral.

Morality and Rationality

Kant claimed that when we act on a maxim that can't be universalized, we are contradicting ourselves. We are being inconsistent. We are assuming that it is acceptable to act in a certain way, even though our purposes could not be achieved if others acted in that very same way. When we

make an exception of ourselves, we are acting as if we were more important than anyone else, and going on as if we were exempt from rules that others must obey. But we are not more important than others, and we are not exempt from these requirements.

It follows that when we behave immorally, we are reasoning badly. We are making mistaken assumptions—that we are more important than other people, that the rules applying to them do not apply to us. Those mistakes, and the inconsistent, contradictory reasoning behind them, show that *immoral conduct is irrational*.

That is a very striking claim, and one that most of us hope is true. We want to be able to convict rapists or terrorists of irrationality, of ignoring their strongest reasons. We want to be able to truthfully say that there were excellent reasons for them to do good and to avoid evil. Kant believed that we could do this.

But how could Kant be right? Consider the ruthless contract killer who knows precisely what he wants, knows exactly how to get it, and executes his plan without fail. Morality doesn't enter into his calculations. He knows that what he is doing is immoral, but that doesn't faze him. It seems that such a person is reasoning flawlessly. How can we defend the claim that he is irrational?

Let's call this the *Amoralist's Challenge*. The **amoralist** is someone who believes in right and wrong but doesn't care about morality at all. The amoralist has the same attitude to moral rules as I do to the rules of professional cricket—yes, they really exist, but they have absolutely no bearing on my life. Obedience to these rules is completely optional. If I am interested in playing the game, then I'll follow the rules. If not, then there is no reason to do so.

The Amoralist's Challenge supports this view in the following way:

1. People have a reason to do something only if doing it will get them what they care about.
2. Doing their moral duty sometimes fails to get people what they care about.
3. Therefore, people sometimes lack any reason to do their moral duty.
4. If people lack any reason to do their moral duty, then violating their moral duty can be perfectly rational.
5. Therefore, it can be perfectly rational for people to violate their moral duty.

The success of this argument would undermine the thought that morality, all by itself, supplies us with good reason to do as it says. It would also refute Kant's claim that immoral actions are always irrational.

Kant thought that you act irrationally when you act contrary to your strongest reasons. And he thought that when moral reasons apply to a given situation, they are always the strongest reasons. Moral reasons are more important than any other kind of consideration. If morality requires you to do something, then that is what you must do—even if you don't want to do it, even if you'll suffer for doing it, and even if the results of doing it cause only misery to others.

Kant admits that the ruthless contract killer, like so many other successful criminals, did, in a sense, reason perfectly well. He followed what Kant called **hypothetical imperatives**. Specifically, these imperatives (commands) are commands of reason. They command us to do whatever is needed in order to get what we care about. Hypothetical imperatives tell us how to achieve our goals. They require us, on pain of irrationality, to do certain things, but only because such actions will get us what we want.

For instance, if my goal is to lose twenty pounds (as it often is), then reason requires me to forgo that pint of luscious coffee ice cream. If I want to get that Wall Street job, then reason requires that I line up a good summer internship. Reason demands that I look both ways at a busy intersection if I want to remain alive. These rational commands apply to me because of what I care about. I am irrational if I disregard them or act in a way that violates them.

But what if I decide that I don't care about weight loss or the benefits that weight loss can bring? What if I don't care about a Wall Street job? What if I want to die rather than live? In that case, I am no longer rationally required to pass on the ice cream, get the internship, or look both ways before crossing the street. These commands of reason are precarious. Their existence depends entirely on what I want. When my desires change, these rational requirements change or disappear.

Many people think that all rational requirements are like this—that they are all hypothetical imperatives. That's precisely what the first premise of the Amoralist's Challenge states: all of our reasons for action depend on what we care about.

Kant saw the implications of this argument very clearly, and knew that he had to challenge that first premise. In his jargon, what we need is to show how there can be such a thing as a **categorical imperative**. This is also a command of reason. But unlike hypothetical imperatives,

categorical imperatives are rational requirements that do *not* depend on what we care about. They are requirements of reason that apply to everyone who possesses reason—in other words, everyone able to reflect on the wisdom of her actions, and able to use such reflections to guide her actions. Categorical imperatives command us to do things whether we want to or not, with the result that if we ignore or disobey them, we are acting contrary to reason (i.e., irrationally).

Kant thought that *all moral duties are categorical imperatives*. They apply to us just because we are rational beings. We must obey them even if we don't want to, and even if moral obedience gets us nothing that we care about.

One lesson Kant took from his thoughts about the golden rule is that the basic rules of morality do not depend on our desires. If they did, then moral rules would fail to apply to everyone, since our desires can differ from person to person. This would make morality too variable, and make it possible for people to escape from their moral duty just by changing what they want. Kant thought that he was defending common sense when he claimed that morality is, in this sense, universal—that everyone who can reason must obey its commands.

If moral duties really are categorical imperatives, then we act rationally when we act morally, and we act irrationally when we act immorally. Is that sort of view defensible? Can we really justify the claim that it is rational for everyone to act morally—even if we know that, for some people, moral conduct will only undermine their goals?

Kant thought he could do this. Consider his *Argument for the Irrationality of Immorality*:

1. If you are rational, then you are consistent.
2. If you are consistent, then you obey the principle of universalizability.
3. If you obey the principle of universalizability, then you act morally.
4. Therefore, if you are rational, then you act morally.
5. Therefore, if you act immorally, then you are irrational.

It does seem that rationality requires consistency, as the first premise asserts. And, as we have discussed, the principle of universalizability is a demand of consistency. So, while more could certainly be said about these first two premises, let us take them for granted here and focus on the third. This is the claim that obedience to the principle of universalizability guarantees that our conduct is moral.

Its location in this argument tells us that the principle of universalizability is a crucial element in Kant's reply to the Amoralist's Challenge. He needs to successfully defend the principle in order to secure the claim that rational people are moral people, and immoral people are irrational. Can he do it?

Assessing the Principle of Universalizability

Unfortunately, the principle of universalizability fails as a general test for the morality of our actions. Look at premise 3 of Kant's Argument for the Irrationality of Immorality. It says that a maxim's universalizability is a guarantee of an action's rightness. That is false. We can act on universalizable maxims and still do wrong.

The principle of universalizability seems to be a very attractive way of pointing out how unfairness and inconsistency lead to immorality. So, for instance, when a thief robs a bank in order to gain riches, Kant can show why the robbery is immoral. If everyone acted on the thief's maxim, there would be no money in the bank to steal, and the thief's goal could not be achieved. But what if the criminal had robbed the bank in order to cripple it and put it out of business? If everyone acted that way, then the thief's goal *could* be achieved. So the principle of universalizability fails to condemn the robbery. And yet such an act is surely wrong.

Suppose that someone wants nothing more than to have a picture-perfect lawn. His basic maxim is to do whatever it takes to preserve the beauty of his grass. This would be strange, but that is no strike against his plans. The real thorn in this guy's side is the mail carrier, who always trots across the lovely yard on his way to deliver the mail. He has been warned, repeatedly, to use the sidewalk—but he never listens. So the homeowner decides to take matters into his own hands. The next time the mailman walks across the lawn, the homeowner pulls out a gun and kills him.

That's clearly the wrong thing to do. But this man's maxim is universalizable. It tells him to take whatever steps are necessary (including killing) in order to preserve the beauty of his lawn. What if everyone were to act upon this maxim? Could the homeowner's goal be met? Undoubtedly. And so, if the principle of universalizability is true, this is no murder; it is justifiable killing. That has to cast serious doubt on the principle.

This unlikely example is just an instance of a general problem, one that we have seen before. Recall the case of the fanatic that came up when

we were discussing the *what if everyone did that?* test. The goals of fanatics are ones that can often be met in a world in which everyone shares their aims. Fanatics (such as the lawn fanatic) need not make exceptions of themselves. The murderous aims of any number of groups could easily be achieved in a world in which everyone supported them. Thus fanatics can be consistent in the relevant sense: their guiding principles could be fulfilled if everyone else were to adopt them.

I think this shows that the principle of universalizability fails to give us an adequate test of fairness, for we can follow its advice while still singling out individuals or groups for discriminatory treatment. There can be consistent Nazis, after all. It doesn't follow that their policies are fair or morally acceptable.

Integrity

While utilitarians think of benevolence (the steady commitment to do good for others) as the central moral virtue, Kant touts integrity. Having integrity is living in harmony with the principles you believe in. It is the virtue of consistency. Integrity requires that you resist making an exception of yourself. It demands that you follow your principles even when doing so comes at a real cost. Kant is surely right that there is something admirable about integrity.

But integrity is not the only moral virtue, and it isn't even the most important one. The example of the fanatic shows this. We may be absolutely dedicated to our principles, but if those principles are deeply flawed, it would be better, morally speaking, for us to have less integrity. That's surely what we want of the Nazi commandant or the dedicated terrorist. It would be better were they less principled people. We want them to be more flexible, and more open to the possibility that their guiding ideals are mistaken. When Huck Finn beats himself up for continuing to hide Jim, the escaped slave who accompanies him down the Mississippi, we applaud Huck's lack of integrity. A Huck with greater integrity is also one who would betray Jim's location. We want Huck to be less than fully conscientious, since that will mean Jim's freedom.

Integrity is worthy of our admiration only when it is tied to morally legitimate principles. The problem, as we have seen, is that people of integrity may still be doing wrong. Refusing to make an exception of myself is no guarantee that my principles are morally acceptable. It's not that consistency is worthless. But it fails as a general test for the morality of the principles we live by.

Kant on Absolute Moral Duties

Kant thought that certain sorts of actions are never permitted. Lying is one of them. In a much-discussed case, that of the inquiring murderer, Kant has us imagine a man bent on killing. This man knocks at your door and asks if you know the location of his intended victim. You do. Should you reveal it? If you do, your information is almost certainly going to lead to murder.

Kant thought you had two decent choices. Ideally, you'd just say nothing. That wouldn't help the murderer, and it wouldn't involve lying. But what if you have to say something? In that case, you have to tell the truth—because you must never lie, under any circumstances.

I think that this is the wrong answer, and the interesting thing is that Kant's own theory does not require him to give it. Kant was so convinced that lying was wrong that he misapplied his own theory.

Kant never provided an argument for the claim that the moral rules that prohibit such things as lying and killing are **absolute** (i.e., never permissibly broken). The closest he came to supplying such an argument was in his belief that moral considerations are more important than anything else. In any conflict between moral duty and other demands—say, those of the law, self-interest, or tradition—morality wins.

Still, it doesn't follow that moral duties are absolute, for even if they always outweigh other kinds of considerations, moral duties might conflict with other moral duties. And if they do, they can't all be absolute. Some of them must give way to others.

And can't moral duties conflict with one another? It seems, for instance, that there is a duty to avoid hurting people's feelings, a duty not to start a panic, and a duty to protect innocent people from dangerous attackers. It also seems that fulfilling each of these duties will sometimes require us to lie, and that there is a moral duty not to do so. Perhaps none of these is really a moral duty. Or perhaps, implausibly, we'd never need to lie in order to respect these duties. But it's much more likely that these are real duties, and that they really can conflict with one another. And if that is so, then these duties cannot all be absolute.

This does not spell disaster for Kant. He does not need to defend the existence of absolute moral duties. His philosophy can, for instance, justify lying to the inquiring murderer. Kant's hatred of lying made him overlook a crucial element of his own view—namely, that the morality of an action depends on one's maxim. He just assumed that anyone who lied would be

operating with a maxim like this: tell a lie so as to gain some benefit. That maxim is not universalizable. In a world in which everyone did this, no one could trust the words of others, and people would be unable to obtain any of the goals they were trying to achieve through lying.

But Kant's maxim is not the only one you could have in such a situation. A maxim is a principle that you give yourself. No one forces it on you. When confronted with a potential killer, I might adopt this maxim: *say whatever I need to say in order to prevent the murder of an innocent person*. That maxim is universalizable. The goal I am aiming for—to save an innocent person's life—could be achieved if everyone acted this way.

For Kant, we can't determine whether an act is right or wrong until we know its maxim. And for any given action, there are countless maxims that might support it. After all, we make up our own maxims, and mine may be very different from yours. It follows that there is only one way for Kant to absolutely ban a type of action. And that is to be sure in advance that, of all the hundreds or thousands of maxims that might support an action, *none* of them is universalizable. It is hard to see how we could ever know that.

As a result, it is much harder than Kant thought to defend the existence of absolute moral duties. And in this particular case, that is all to the good, since it opens up the possibility that it is sometimes acceptable to lie—for instance, to the inquiring murderer. Of course, if Kant is right, then we would have to have a universalizable maxim that permits this. But nothing Kant ever said should make us think that this is impossible. Contrary to Kant's personal view, we don't have to regard all (or perhaps any) moral duties as absolute.

Discussion Questions

1. Explain the difference between the golden rule and the *what if everyone did that?* test. What problems arise for each? Do you think that they can be remedied?
2. What is a maxim, and what does it mean for a maxim to be *universalizable*? Why does the principle of universalizability fail to be a good test of the morality of our actions?
3. According to Kant, it is always irrational to act immorally. What reasons does he give for thinking this? Do you agree with him?
4. What is the difference between hypothetical and categorical imperatives? Why did Kant think that morality consists of categorical imperatives?

5. Why does the existence of fanatics pose a challenge to Kant's moral theory? How do you think that the Kantian should respond to this challenge?
6. Is integrity always a virtue? Why or why not?
7. Explain how one would show that a moral duty is absolute according to Kant's moral theory.