

Consequentialism

Francesca Lagioia

CONSEQUENTIALISM

The morality of an action is judged solely by its consequences.



- Focus on the End Result
- Maximize Overall Well-being
- Utilitarianism is a key form

An action is morally required

- iff it delivers that best outcome, relative to its alternative
- iff its good outcomes outweigh its negative outcomes to the largest extent
- iff it produces the highest utility?

Morality as an **optimisation problem** to ensure good outcomes significantly outweigh negative ones

How to measure and compare different types of utility to determine the best course of action?

- What are the good and bad things to be maximised?
- How many there are?
- How much each of them matters?
- Can we construct a single utility function that combines gains and losses over multiple valuable goals?

The reference approach:

Utilitarianism



Core Idea

Primarily developed by **Jeremy Bentham** (1748–1832), utilitarianism states that the **best action** is the one that produces the **greatest happiness for the greatest number**. It focuses on maximizing overall well-being.

Principle of Utility

John **Stuart Mill** refined this in "Utilitarianism" (1861). Actions are **right in proportion as they promote happiness and wrong as they produce the reverse**. Happiness means **pleasure** and **absence of pain**.

Maximizing Well-being

The principle of utility involves **maximizing the overall amount of well-being**, considering everyone affected by our actions.

- **Intentions are irrelevant;**
- **Only results matter** in utilitarianism.

The utility of everybody has to be taken into account equally

Utilitarianism: Examples



Lying to Avoid Harm

Scenario: A doctor lies to a patient to prevent severe distress or save their life.

Consequentialist View: Lying is acceptable if it results in a positive outcome, easing suffering or aiding recovery.



Sacrificing One to Save Five

Scenario: Sacrificing one life to save five others.

Consequentialist View: Actively causing harm to one is justified if it saves multiple lives, reducing overall suffering.




Stealing to Feed a Family

Scenario: A desperate parent steals food to feed starving children.

Consequentialist View: If the benefit of saving children outweighs the harm of stealing, the act is morally justified.



Utilitarianism – Concrete Examples in AI



Chatbot Lies “for User”

Scenario: A depressed user asks a chatbot if his/her situation will improve.

Consequentialist View: The chatbot might offer a “**white lie**” to provide comfort, if that lie leads to a better emotional outcome.



Autonomous Car Decision

Scenario: An autonomous car must choose between hitting 5 pedestrians or swerving to hit 1. **Consequentialist View:** AI chooses to swerve and hit 1, reducing overall death, even if it involves harm.

Advantages of utilitarianism

- **Conceptually simple:** easy to understand and apply, grounded in the doctrine of impartiality
- **Egalitarian:** everybody's utility counts in the same way
- **Fits with some basic intuitions:** making people happy is good, making them suffer is bad, reinforcing its common-sense appeal.

Example(s): Slavery, rape, humiliating defenseless people, and killing innocent victims are condemned, aligning with our moral judgments. Helping the poor, keeping promises, telling the truth, and facing danger are commended, which also resonates with our moral beliefs



Advantages of utilitarianism

- **Moral Conflict Resolution:** offers concrete guidance on how to solve conflicts due to its **single ultimate rule**: maximize well-being
- **Flexibility:** great flexibility, as no moral rule is absolute except the principle of utility. Even rules against cannibalism or torture can be violated if overall well-being rises. Adaptability allows for **context-specific ethical decisions**, challenging rigid moral codes and fostering a **dynamic approach to complex moral issues**. This ensures **decisions are tailored to the specific circumstances**.



Two Versions of Utilitarianism: Act vs. Rule

Act Utilitarianism Definition

Assesses each individual action based solely on **the specific consequences that action produces, maximizing net benefit.**

Key Idea

The rightness of an action is determined on **a case-by-case basis.** Even if an action usually considered wrong might lead to a better outcome in a unique circumstance, it can be justified.

Example

Imagine lying to spare feelings or breaking a rule in a crisis.
Act utilitarianism might justify these actions if they yield a better outcome in these specific cases.

Two Versions of Utilitarianism: Act vs. Rule

Rule Utilitarianism Definition

Rule utilitarianism evaluates actions based on whether they conform to **rules** that

- **if generally followed, would lead to the greatest overall happiness.**

Key Idea


Instead of assessing each action individually, rule utilitarianism endorses **following established rules** for better **long-term results** and social stability.

Examples

Truth-telling: Even if lying might yield immediate benefits, the general rule against lying is crucial for trust. Traffic laws are also vital for ensuring overall safety.

Act vs. Rule Utilitarianism

Versions	Act Utilitarianism	Rule Utilitarianism
Focus	Specific consequences of each individual act	Overall benefits of adhering to a set of rules
Decision Process	Case-by-case evaluation	Consistency with established rules
Example - Lying	Lying to comfort a friend in a unique situation	Upholding a rule of truth-telling to maintain trust
Example – Traffic Laws	Breaking a law in an emergency if it saves a life	Following traffic laws to ensure overall safety



Issues with act utilitarianism

➤ Decision Procedure

- Does act utilitarianism provide a good decision procedure?
- Can we choose actions by optimizing outcomes?
- Do we have enough information for this calculation, or can AI systems have enough?

➤ Standard: Does it provide a good standard for assessing decisions?

- Act utilitarianism demands a case-by-case analysis of outcomes, which can be highly impractical (?).
- It's often impossible to accurately predict all consequences.

Example: A business deciding whether to lay off employees to increase profits must consider long-term social impacts, making ethical decision-making challenging



Issues with act utilitarianism

➤ Questionable Acts?

- Focusing on immediate consequences, can its moral flexibility lead to questionable acts, conflicting with intuitions or rights?

Example: Torturing a suspect to prevent a terrorist attack, seen as inherently wrong, might be justified if it saves many lives. This clashes with basic moral



Issues with act utilitarianism



➤ Too demanding and Impartiality problem?

- it forces individuals to consider the utility impact of every single action and be impartial

Example:

- Should I give the same importance to everybody, regardless of their connection to me? Morality seems to recommend partiality to family and friends. E.g., It seems right to care about your children more than others, but utilitarianism rejects that a person, just because he/she is your son or friend, is more deserving.
- Is it OK to harm some people for the greater benefit of others? Reprisals? Torture? Sadism?



Issues with act utilitarianism



➤ Constant self-sacrifice?

- Anytime you can do more good for others than for yourself, you are required to do so. This implies foregoing personal pleasures for the greater good.

Example:

- Choosing between a vacation and providing basic water kits for 100 families highlights the constant self-sacrifice demanded by strict utilitarianism.



Issues with act utilitarianism




➤ Moral overload?

- The constant utility calculation might be overwhelming in everyday decision-making

Example:

- A person evaluating every small decision by calculating its impact on happiness can experience decision paralysis and moral burnout, showing act utilitarianism's



Rule Utilitarianism

Handling Objections

Rule utilitarianism promises to **handle a number of objections to act utilitarianism** by focusing on **rules** rather than individual actions.

Optimific Social Rule

An action is morally right because **it is required by an optimific social rule**, i.e., a **rule that would be beneficial if nearly everyone followed it**.

Basic Idea

Rather than determining an action's morality by asking about its results, **we ask whether the action conforms to a moral rule**.

- It's generally okay to tell the truth and not steal because general compliance with such norms would deliver **the greatest utility**.
- Rule utilitarianism aims to balance individual actions with overall societal well-being



Rule Utilitarianism

1

Describe the Rule

Carefully describe **the rule to be evaluated**. Understanding the specific content of the rule is essential for assessing its potential impact.

2

Imagine Society

Imagine what a society would be like if just about **everyone in it endorsed the rule**. This step is crucial for **predicting broad societal consequences**.

3

Ask the Question

Will that society be better off with this rule than with any competing rule? This question determines whether the rule is **truly optimific**.

4

Optimific Social Rule?

If the answer to the question is **yes**, then this rule is an **optimific social rule**. If the answer is no, then it isn't a genuine moral rule.

Rule Utilitarianism: Open Questions

Exceptional Cases

What about those **exceptional cases in which the rule does not deliver?** How should we handle **situations** where following the rule leads to suboptimal outcomes?

Non-Compliance

What if you know that **most people are not following the rule?** How does widespread non-compliance affect the moral imperative to adhere to the rule?

Honesty Dilemma

Should we be honest if most people around us are dishonest? How should we act when the prevailing social environment deviates from the optimific rule?



A Further Issue: Distribution



Distribution Matters?

- Does it matter **how the good and bad outcomes are distributed**?
- Is it okay to take an action that benefits some to the detriment of others if overall benefits outweigh disadvantages?

Wealth Redistribution

Utilitarianism favors (modest) redistribution of wealth, since **the same amount of money gives more utility to the poor than to the rich**, but wealth generation must be considered.

Wealth Maximization

Wealth maximization aims at maximizing wealth in society **regardless of distribution**, adopted by some economic approaches. This contrasts with utilitarianism's emphasis on well-being.

The Trolley Problem



Ethical Thought Experiment

The trolley problem is a famous ethical thought experiment in moral philosophy, particularly relevant in **utilitarianism** and **deontology**, presenting a **moral dilemma**.

Classic Scenario

A trolley is heading down a track and will kill 5 people. You can switch it to another track where it would kill 1 person.


The dilemma is whether to act or not

- ☐ Is it morally ok to actively cause harm (pull the lever) to save more people, or is it better to do nothing and allow more deaths to occur?

Options

- **Option 1:** Do nothing, and the trolley kills 5 people.
- **Option 2:** Pull the lever, diverting the trolley to kill 1 person instead.

- ☐ **Which action is more ethical?**
- ☐ **What should AI tasked with monitoring traffic do? How should AI be programmed?**




Utilitarianism

Utilitarianism focuses on the **outcome**, asserting the morally right action is the one that **maximizes overall well-being** by **minimizing harm** and **maximizing benefit**.


Pulling the Lever

A utilitarian would argue that **pulling the lever** is the right choice since it minimizes total harm, aligning with consequentialist principles.



Different Trolley Paths


Visualizing **various trolley paths** helps illustrate the complexity of the ethical problem. **The different paths can lead to various outcomes, each with distinct moral implications**, prompting critical analysis and ethical evaluation.



the switch
Foot, 1967



the fat man
Thomson, 1976



the fat villain



the loop
Costa, 1987



the man in the yard
Unger, 1992

Ethical Perspectives on the Trolley Problem



The Fat Man Variant

Instead of a lever, you are on a bridge with a very **fat person**. Pushing them onto the tracks will stop the trolley and save five people, but they will die.




Direct Physical Involvement

This scenario tests whether **direct physical involvement changes the morality of the decision**, as **pushing** someone not at risk feels different from **pulling a lever** to decide who should survive among those at risk.



The Ethical Dilemma of Autonomous Vehicles



Bonnefon et al. 2016


This study explores the **ethical considerations** in **programming autonomous vehicles**, particularly focusing on how they should handle **unavoidable accident scenarios**.

Programming Decisions


How should Avs be programmed to prioritize safety in scenarios where harm is inevitable? Should they minimize overall harm or protect occupants?

Start Judging

What should the self-driving car do?



Show Description



Show Description

<https://www.moralmachine.net>

Judith Jarvis Thomson: The Surgeon Case



The Scenario

A surgeon has **5 patients** needing organ transplants but no organs are available. A **healthy traveler** comes in for a checkup and is a perfect match for all five.

Is it moral to **sacrifice 1 life to save multiple lives?** The dilemma involves balancing the **utilitarian goal of maximizing lives saved** with the **ethical prohibition against directly causing harm**.



Ethical Dilemma

If the young man disappear, no one would suspect the doctor. **Is it morally justifiable to kill him and use his organs to save the 5 dying patients?**



Questions?

