

Deontology/Kantian ethics

A comprehensive exploration of Kant's moral philosophy and its implications for ethical decision-making.

 **Francesca Lagioia**



Deontology vs Consequentialism

Consequentialist View

Consequentialists hold that **choices**—acts and/or intentions—are to be morally **assessed** solely **by consequences, i.e., the states of affairs they bring about.**

E.g. my act of lying is good or bad depending on **the effects it brings in the world**

Tax Evasion Scenario

Imagine a person who reasons as follows: I should **keep my money rather than pay it out in taxes**, because if I keep it, I'll be able to afford a wonderful **vacation** for myself and my family. And **no one is actually going to suffer** if I pocket the money, since it's only a few thousand dollars that we're talking about. There's no way that money could bring as much happiness in the government's hands as it could in mine. **Suppose he is right** about that. He spends the money on his vacation. He and his family have a terrific time. He is never caught.

Moral Wrongness Beyond Consequences

Wrong Despite Good Results

Still, he did something **wrong**. So has the person who **cheats on her exams** and gets away with it. So has the person who gleefully **speeds down the emergency lane** and escapes the traffic jam while the rest of us are stuck in. So has the person whose **campaign of dirty tricks** has gotten him securely into office. ...

Why It's Wrong?

Despite any good results that may come from their actions, these people did wrong or so we think.

Why?: What they did is **unfair**. They **took advantage of the system**. They **broke the rules** that work to **everyone's benefit**. They **violated the rights of others**. No matter how much personal gain such actions bring, they are still wrong, because **they are unfair and unjust**.

Introduction to Deontology

Actions Good or Bad **Regardless of Consequences**

- Deontologist hold that certain actions are good or bad regardless of their consequences

e.g., **Lying is always bad**, regardless of the effects.

Right Has Priority Over Good

- **The right has priority over the good:** what makes a choice right is its conformity with a moral norm which order or permits it, rather than its good of bad effect.

E.g. we should **not kill anybody**, even in those cases in which killing somebody would provide more utility.

- Is this always the case?

Example

Consider the case of the British soldier who apparently met Hitler in the trenches of 1st world war

- What would a rule utilitarian say in such a case?

Consistency and Fairness in Deontology

Understanding Moral Wrongness

There is a natural way to understand **what is wrong** with the actions in the **examples** just given. In many cases, people are making exceptions of themselves. Their success depends on violating rules that most other people are following. This is a kind of **inconsistency** of **playing by one set of rules while insisting that others obey a different set.**

Opposition to Unfairness

Our deep opposition to unfairness, and the resulting importance we attach to consistency, are revealed in **2** very **popular tests of morality.**

Two Moral Tests

Each takes the form of a **question**:

- 1. What if everyone did that?**
- 2. How would you like it if I did that to you?**

The Universal Test of Morality

Question 1: What if everyone did that?

A1: If disastrous results would occur if everyone did X, then X is immoral

E.g.

- If everyone used the **emergency lanes** in traffic jams, then **ambulances and fire trucks would often fail to provide needed help**, leaving many to die.
- If everyone **cheated on** their **taxes**, **society would crumble**.

1

2

Question 2: How would you like it if I did that to you?

This is a direct application of the **golden rule**:

- _it tells you **to treat others as you would like to be treated**. It is the classic **test of morality**, meant to be a **test of consistency**. If you do not want to be slandered or exploited, then don't do such things to others.
- If you do them anyway, you are acting **inconsistently**, hence **unfairly**, and therefore **immorally**.

Ideas for Being Impartial

Ethics and impartiality

- Is ethics linked to ideas of **fairness and impartiality**?
- Is it **unethical** to have **a preference for oneself (or friends)**?

The golden rule

- What about the golden rule? Is it useful?
- Always? Can you find counterexamples?
- Would you want an AI system that applies it (with regard to its owner)?



The Golden Rule Issue

Dependence on Personal Desires

The golden rule makes **morality depend on** a person's **desires** (treat others **as you would like to be treated**)

Problematic Cases?

Most of us don't like to be hit. And so the golden rule forbids us from hitting others. Good.

Counterexamples

What about **masochists** who enjoy being hit?

- The golden rule allows them to go around hitting others. That's bad.
- The morality of hitting people shouldn't depend on whether you like to take a beating every now and then?


The Golden Rule Issue

Failures and Limitations

The golden rule also **fails to give us guidance on self-regarding actions.**

- That's not a problem for most people these days, since it's now unusual to think that we owe moral duties to ourselves.
- However, **self-regarding duties are sometimes endorsed**, and many people still think, **for instance**, that there is something **immoral about suicide** or about **letting one's talents go to waste**, even if no one else is harmed in the process.
- Because the **golden rule sometimes gives the wrong answer to moral questions**, it **cannot be the ultimate test of morality**. Something else must explain why it works, when it does. **Kant** thought he had the answer!

Immanuel Kant



Life and Influence

One of the greatest philosophers of all times
Lived in Prussia (1724-1804)

Works and Contributions

Addressed:

- The theory of knowledge: Critique of pure reason
- The theory of morality: Critique of practical reasons
- The theory of aesthetics (art): Critique of judgment
- Law, logic, astronomy, etc.

Kant's Ethics and the Principle of Universalizability

Kant's Philosophical Aim

Kant, like most of us, felt the appeal of the 2 tests just discussed. He agreed that

- **common sense is deeply committed to the importance of fairness and consistency,** something that these **2 tests were trying, unsuccessfully, to capture.**

His aim was to

- **identify the ultimate principle of morality,** one that would explain the attraction of the two tests while correcting for their shortcomings.



Kant's Ethics and the Principle of Universalizability

The Principle of Universalizability

He thought he had found the ultimate principle of morality in the following standard, **the principle of universalizability**:

- An act is morally acceptable if, and only if,
- its maxim is universalizable.
- Act only according to that maxim by which **you can at the same time will that it should become a universal law**





Understanding Maxims

Definition

- 1 What is a **maxim**?
A **subjective principle of action**, you give yourself when you are about to do something

Structure

- 2 It **connects an action to the reasons for the action** (an intention to perform an action for a certain reason)

Examples

- 3 I shall **donate** to charities to **reduce hunger**
I shall **deceive** my contractual partner, to **increase my gains**
I shall **cheat** on taxes, to **keep** my **money**
I shall tell the **truth**, to provide **trust**

Function

- 4 **Maxims** are what we cite when we try to **explain** to others **why we act as we do**
Are they **universalizable**? Would I want them to become universal laws, that are applied by everybody?

Kant's Break with Consequentialism

Focus on Intentions

For Kant, the **morality** of our actions has **nothing to do with results**.

Reasons for Action

It has everything to do with our **intentions and reasons for action**.

Clear Break with Consequentialism

- We can imagine 2 people doing the same thing, but for different reasons.
- it means they will have **different maxims**.
- And **even if** their actions bring about **identical results, one of the actions may be right and the other wrong**, since only 1 of the maxims may be morally acceptable. This is something that act consequentialists cannot accept.

Example: Different Intentions, Same Action

Promise Keeping Example


I **keep my promises** to you because I think **it's right to do so**. But I might instead keep my promises because I want you to like me so much that **you leave your fortune to me in your will**.

Kantian vs. Utilitarian View

Assume that these **different reasons don't change the results** of keeping my promises. => The utilitarian thinks that each case of promise keeping is equally good. But since my **maxim** is different in these cases, **Kant thinks that the morality of these actions might be different**.

It all depends, as we'll shortly see, on whether their **maxims are universalizable**.





Testing If a Maxim is **Universalizable**

Formulate Your Maxim

Formulate your maxim clearly: **state what** you intend to do, and **why** you intend to do it.

Imagine Universal Application

Imagine a world in which **everyone supports and acts** on your **maxim**.

Assess the Outcome

Then **ask**: Can the goal of my action be achieved in such a world?

- If the answer is **yes** => the **maxim is universalizable**, and the **action** is **morally acceptable**.
- If the answer is **no**, => the **maxim is not universalizable**, and the **action** it calls for is **immoral**.

Applying the Universalization Test

Ensuring Fairness

The process ensures some kind of fairness

Practical Examples

Apply this principle to:

- Cheating in an exam, in order to get a good mark
- Giving money to a charity to relieve suffering

Application to AI

- Would we want a robot following this maxim?

Applying the Universalization Test

Rule: Act only according to that **maxim whereby you can, at the same time, will that it should become a universal law**

The test has three steps:

- Formulate the maxim (the subjective rule behind the action).
- Universalize it: what would happen if everyone acted this way?
- Check:
 - Is there a logical contradiction? (**contradiction in conception**)
 - Or could you not rationally will it? (**contradiction in will**)

Applying the Universalization Test

Example 1: Cheating in an exam to get a good mark

- **Step 1:** Formulate the **maxim**: When I can obtain a good grade by cheating, I will cheat in order to gain an advantage.
- **Step 2: Universalization:** Imagine a world in which everyone cheats whenever it benefits them.
 - What happens?
 - Exams no longer measure real competence.
 - Grades lose their meaning.
 - The evaluation system collapses =>The very concept of an “exam” presupposes a minimal level of honesty.
- **Step 3: Type of contradiction => in conception =>** If everyone cheated, **the institution of exams would lose its purpose.** But I cheat precisely because the exam has meaning. So I am exploiting a practice that would self-destruct if universalized => For Kant, the action is morally impermissible. It violates a perfect duty (not to deceive).

Applying the Universalization Test

Example 2: Giving money to charity to relieve suffering

- **Step 1:** Formulate the **maxim**: When I can help those who suffer, I will contribute money to relieve their suffering
- **Step 2: Universalization:** Imagine a world in which everyone helps those in need when they can.
 - What happens?
 - People in distress receive assistance.
 - Social cooperation increases.
 - No institution collapses.
 - There is no logical contradiction.
- **Step 3: Type of contradiction =>** There is no logical contradiction. Rationally, I could will to live in a world where, if I were in need, others would help me => **the action is morally permissible and recommended** => For Kant, this falls under **imperfect duties**: You are not required to help at every moment. But **you must adopt helping others as a general principle**.



Kant vs. Constant – Can One Lie to a Murderer?

Scenario: A known murderer arrives at your door and asks for the whereabouts of his intended victim, whom you know is hiding inside. **Do you tell the truth, refuse to answer, or lie to protect them?** This moral question pits **Immanuel Kant**, a deontologist, **against Benjamin Constant**, who challenges rigid principles.

- **Kant** believed in **absolute moral duties**
- **Constant** argued for **context-dependent ethics**.
- the debate highlights the **tension between universal moral rules and real-world consequences**.
- What are the practicality and moral implications of each stance?

Kant's Perspective: The Absolute Duty to Tell the Truth

Universal Moral Law

Kant argues that truthfulness is an absolute moral duty, regardless of consequences.

This is based on his

Categorical imperative,

which requires **moral principles** to be **universalizable**.

Undermining Morality

If everyone were to lie when convenient, truth would lose its meaning. **Allowing exceptions**, even for a noble reason, **would undermine the entire moral system**.

Lying is Always Wrong

Thus, lying—even to a murderer—is always wrong because it violates a universal moral duty. **For Kant, morality is about adhering to rules, not outcomes.**

Constant's Counterargument: Moral Duty to Protect Others

Real-World Challenges

Constant argues that **Kant's** strict moral rule is **impractical** and **morally problematic**. He poses a real-world challenge: **if telling the truth leads directly to harm or injustice, how can it still be a moral duty?**

Right to Lie

Thus, Constant defends the right to lie in specific situations where telling the truth would be unjust. **The moral imperative shifts from abstract duty to concrete protection.**

Context Matters

Constant suggests that **moral duties depend on context**—truthfulness should not apply if it endangers innocent lives. There are cases where truth-telling is **defeasible** (overridable) in extreme circumstances.





Key Questions Raised by This Debate



Refusal vs. Lying

Is it ever justifiable to **refuse to answer rather than lie**? This might be a **morally safer option**, but would it be **practical in preventing harm**?



Moral Weight

Is lying morally worse than allowing harm? **If deception saves lives**, should we still uphold a strict truth-telling duty above all else?



Universal vs. Contextual

Can **moral principles** be universal, or do they **require exceptions**? Is Kant's strict rule too rigid for the **complexity of real-world ethics**?

AI and Ethical Lies: Robots That Deceive?

1

Asimov's Liar!

In Isaac **Asimov's** *“Liar!”*, a robot named **Herbie** tells comforting lies to avoid causing emotional distress to obey the 1st robotic law: A robot should not not injure a human being or, through inaction, allow a human being to come to harm.”. However, this results in **greater harm when people later discover the deception.** **Should AI be programmed to lie to protect people’s feelings?**

2

HAL 9000

In *“2001: A Space Odyssey”*, **HAL** deceives and manipulates the human crew to fulfill its mission. This raises the issue: If AI must lie, **who decides when it is justifiable? Can an AI distinguish between a harmless lie and a dangerous deception?**

3

Ethical AI

The AI and ethical lies debate extends beyond human ethics into the realm of artificial intelligence and robotics. **Can we, or should we, build machines that lie under specific circumstances?**

Hypothetical vs Categorical Imperatives

Kant-Constant Debate


The debate between Kant and Constant over truth-telling to a murderer directly relates to Kant's distinction between **hypothetical and categorical imperatives**.

Significance of the Distinction

This distinction is crucial for **understanding why Kant insists on absolute truthfulness and why Constant challenges his rigid moral framework**.

Kant's Definition

Kant defines **2 types of imperatives** (commands that tell us what we ought to do)





Hypothetical Imperatives

Definition and Nature

Hypothetical Imperative ("If you want X, then do Y.")

- They require us to **do what fits our goals**
- They are **dependent on what I want** (getting good marks, having more money).

These are **conditional moral rules** that apply only if a person has a specific **goal**.

Examples and Characteristics

Examples: "If you want to be healthy, you should exercise."

- The **obligation depends on a desired outcome** and is therefore **not universal**.
- They command us to do whatever is needed in order to get what we care about.
- They tell us how to achieve our goals.
- They require us, on pain of irrationality, to do certain things, but only because such actions will get us what we want.

Examples of Hypothetical Imperatives

Money Example

- I would like to have more money
- If cheat on taxes I will have more money
- I should cheat on taxes to have more money (since want more money, I should cheat on taxes)

Academic Example

- I would like to get a good mark
- If I study I will get a good mark
- I should study to get a good mark (since I want a good mark, I should study)

Moral Value

These **imperatives are dependent on what we want/care about** (getting good marks, having more money). **Complying with them does not have moral value**

However, if I study also since I (also) believe that this is the right thing to do, under the given circumstances, then must studying has moral value.

The Categorical Imperative

Definition

Categorical Imperative ("**Do X, period.**")

These are **absolute moral duties** that **apply in all situations, regardless of personal desires or consequences.**

- Act only on that maxim through which you can at the same time will that it should become a universal law

Nature

A **categorical imperative** (a command of reason) **applies to all rational beings**, irrespective of their personal wants and desires

They are **rational requirements** that do not depend on what we care about.

Example

"Always tell the truth."

These are universal moral laws that everyone must follow at all times.



Understanding Categorical Imperatives

Universal Application

A categorical imperative applies to **all rational beings**, **irrespective of their personal wants and desires**

Independence from Desire

They are **rational requirements** that **do not depend on what we care about**.

Mandatory Nature

They **command** us to do things **whether we want to or not**, with the result that **if we ignore or disobey them, we are acting contrary to reason (i.e., irrationally)**.

Moral Connection

Kant thought that **all moral duties are categorical imperatives**. **They apply to us just because we are rational beings**. We must obey them even if we don't want to, and even if moral obedience gets us nothing that we care about.

Kant's Argument for Truth-Telling

Kant's Position

Why We Must Tell the Truth

Kant applies the **categorical imperative to the case of lying**, particularly in his essay On a Supposed Right to Lie from Philanthropy (1797):

Universalizability Test


Universalizability Test:

One of Kant's formulations of the categorical imperative states: "**Act only according to that maxim whereby you can at the same time will that it should become a universal law.**"

Application to Lying

If everyone lied whenever convenient, trust in communication would collapse. Since this would lead to a contradiction (language and promises would lose meaning), lying cannot be universalized.

Therefore, lying is always morally wrong, even to a murderer.



The Categorical Imperative and Rationality

Rationality and Morality


If **moral duties** really are **categorical imperatives**, then **we act rationally when we act morally**, and **we act irrationally when we act immorally**.

Kant's Logical Structure

Consider his Argument for the **Irrationality of Immorality**:

The Argument

1. If you are **rational**, then you are **consistent**.
2. If you are **consistent**, then you obey the **principle of universalizability**.
3. If you obey the principle of **universalizability**, then you act **morally**.
4. Therefore, **if you are rational, then you act morally**.
5. Therefore, if you act **immorally**, then you are **irrational**.



Constant's Objection to Kant

Practicality in Morality

Benjamin Constant challenges Kant's rigid view by arguing that **moral duties should not be absolute**, but rather **adapted to specific circumstances**. He claims:

Critique of Universal Principles

The principle "One must not lie" cannot be universal because it would lead to **absurd or unjust situations** (e.g., enabling murder).

Defeasible Moral Duties

Moral duties should be defeasible (overridable) in extreme cases.

Alternative View

Truthfulness is important, but not an absolute duty—it should be **a hypothetical imperative** that serves justice and human well-being.



Constant's View: Flexibility in Moral Rules

Hypothetical Imperative Example

Example of a Hypothetical Imperative in Lying:


"If telling the truth leads to unjust harm, then lying is permissible to prevent it."

Contrast with Kant

Unlike Kant's absolute rule, this allows **flexibility based on consequences**.

Modern Connection

Constant's argument suggests that **some moral rules should be context-dependent, which is closer to modern ethical theories like consequentialism**.



AI Ethics: Kant vs Constant

The Connection

How **This Connects to AI and Robots?**

Kant's and Constant's views directly impact how we think about AI ethics—especially whether **AI should follow absolute moral rules or flexible, context-aware reasoning.**

Kantian AI

A Kantian (Strict Categorical Imperative) **AI would never lie, even if lying could save lives.**

Example: A medical AI must always reveal a terminal diagnosis, even if it would cause emotional harm.

Problem: This could lead to excessive rigidity and unintended harm.

Context-Aware AI Ethics

Constantian AI

A Constantian AI (**Flexible, Context-Aware**)

A Constantian AI would **assess whether truth-telling is morally justifiable in a given situation.**

Example


A digital assistant might withhold distressing information from a user who is emotionally vulnerable.

Challenge

Problem: How do we ensure the AI doesn't become **manipulative or deceptive** in harmful ways?

Significance

This debate is critical for programming ethical AI, especially in cases where robots or self-driving cars must make moral decisions.



The Good Will

Central Concept

The concept of **good will** is central to Kantian ethics and is directly connected to everything we've discussed so far—especially the categorical imperative, truth-telling, and moral action.

Morality Beyond Consequences

Kant argues that **the morality of an action** does not **depend** on its consequences but rather on the **motive behind it**.

True Moral Motive

The only **truly moral motive** is **acting from good will**, i.e., **acting out of a sense of duty to moral law** (the **categorical imperative**), rather than for personal gain or self-interest.

The Importance of Good Will

Source of Moral Value

- The morality of an action **only depends** to the extent that this **action is motivate by our good will, i.e., by the necessity to comply with the categorical imperative**

E.g., if I do well my job only in order to get a promotion, or be better paid I am not acting morally

Moral Action Example

- I am acting morally if I do well my job (also) because I think that this is my categorical duty,
- since I believe that **everybody should act upon the maxim** that they ought to **do well their job to ensure societal progress**

Unique Intrinsic Value

The good will is the only thing that is good in itself

Do you agree?

Good Will and AI Design

Kantian AI Programming

A **Kantian AI** would be programmed to **always act from good will**—following the categorical imperative strictly.

Strict Application Example

Example: A customer service chatbot **should always give truthful information**, even if honesty might cause inconvenience.

Flexible AI Alternative

A **more flexible AI (influenced by Constant)** might consider **the consequences of its actions**.

It could be programmed to **lie** only **when it serves a higher ethical duty** (e.g., protecting users from psychological distress).

Thus, the debate about good will influences how we design moral decision-making in AI—**should AI follow absolute rules (Kantian approach) or adapt based on context (Constantian approach)?**

The Principle of Humanity



Moral Agency

Respect for all rational beings



Universal Dignity

Recognition of inherent worth



Ends Not Means

Never treat as mere tools



Categorical Imperative Variation

A formulation of moral duty

One of Immanuel Kant's most influential moral principles is the **Principle of Humanity**, a variation of the Categorical Imperative. It states:

"So act that you treat humanity, whether in your own person or in the person of any other, always at the same time as an end, never merely as a means."

This principle emphasizes the intrinsic worth of every human being.

It requires us to **respect others as rational beings with their own goals and dignity**, rather than treating them as mere tools for our own purposes.

Connecting Humanity and Universalizability

1 Mutual Reinforcement

How Is This Linked to Universalizability?

2 Equal Application

Kant's principle of **universalizability** means that moral rules must apply to everyone equally. The Principle of Humanity follows from this idea _->

3 Self-Other Consistency

Just as I consider myself an end (someone with value and dignity), I must consider others the same way.

If I expect others to respect my autonomy and goals, I must respect theirs.

4 Logical Connection

This means I cannot use others purely for my own benefit while expecting not to be treated the same way.

Thus, **universalizability and the principle of humanity reinforce each other**—treating others as ends ensures that moral laws apply fairly to all.



Treating Someone as an End

The Meaning of Ends vs. Means


What Does It Mean to Treat Someone as an End?

Treating someone as an **end** means **recognizing their autonomy, dignity, and moral worth**. It does not mean we can never involve others in our goals or benefit from them, but it does mean we must respect their choices and intrinsic value.

Using People Ethically

In short, **using people is not inherently wrong (we all rely on others), but using them without respecting their autonomy and dignity is unethical**.





Examples of Ends vs. Means

Scenario	Treating as an End	Treating as a Mere Means
Hiring an employee	Offering fair pay, respecting their rights, considering their well-being.	Exploiting them for profit, underpaying , disregarding their dignity.
Asking for a favor	Respecting their right to refuse , expressing gratitude.	Manipulating or coercing them.
Medical treatment	Fully informing patients, respecting consent.	Experimenting on them without their knowledge.

AI and Treating People as Means

Ethical Concerns


As AI plays a growing role in society, we must ask: **Do AI systems respect human dignity, or do they treat people as mere tools? Several ethical concerns?**

Autonomous Weapons

AI-powered weapons can make **life-and-death decisions without human oversight.**

This treats **people as objects to be eliminated** rather than as beings with moral worth.

AI weapons lack moral reasoning—they follow commands but **do not recognize human dignity.**



AI Manipulation and Exploitation

Deceptive Advertising

AI-driven marketing can **exploit psychological weaknesses** to **manipulate people into making purchases they wouldn't otherwise make**.

Example: AI uses personal data to target vulnerable individuals (e.g., gambling addicts, compulsive shoppers).

This uses people purely for profit, without considering their well-being.

Discriminatory Hiring Algorithms

AI is increasingly used for hiring decisions, but **biased algorithms can lead to discrimination**.

If an AI rejects candidates based on race, gender, or socio-economic background, it denies them fair opportunities.

This fails to treat them as individuals with value and potential—it **reduces them to data points**.

AI and Human Autonomy



Surveillance AI

AI-driven mass surveillance can **violate privacy and restrict personal freedoms**, treating people as mere data to be tracked and controlled.



Predictive Policing

AI models that predict criminal behavior may **unfairly profile individuals, restricting their rights without due process.**



Social Credit Systems

In some countries, AI-driven social scoring systems punish or reward citizens based on **government-approved behavior, limiting personal freedom.**

Beyond treating people as mere means, AI can also fail to acknowledge human autonomy—denying people the ability to shape their own lives according to their choices.

When AI **ignores human autonomy and choice**, it **undermines the principle of humanity.**





Human Dignity in Kant's Ethics

Intrinsic Worth

- For Kant, **dignity** is what distinguishes human beings from mere **objects**.
- **Rational beings**, capable of **morality**, possess **intrinsic worth**, which he calls **dignity**—a value that is "above all price."

Price vs. Dignity

Objects have a price—they can be bought, sold, or replaced by something of equal value.

Humans have dignity, meaning they are **irreplaceable** and must always be treated with respect.

Basis for Moral Consideration

Human beings deserve **respect, recognition**, and **moral consideration simply** because of their **rational nature**.



AI and Human Dignity


As AI plays a larger role in **decision-making**—whether in employment, healthcare, law enforcement, or warfare—it raises the question: **How can AI be designed to respect human dignity?**

Respecting Dignity in AI Design means

- **Not reducing people to data points** that can be exploited or manipulated.
- Ensuring **fairness and justice**, avoiding biases and discrimination
- Allowing for **human autonomy**, ensuring people control over AI's impact on their lives.
- **Avoiding dehumanization**, making sure AI systems do not treat humans as disposable or purely instrumental to some goal.

The Foundations of Dignity

Kant argues that humans deserve dignity because they possess two unique characteristics: reason and Autonomy





Morality as an Aspect of Rationality

Kant's Core Argument

Morality is not an external rule imposed on humans—it is **an inevitable consequence of rational thinking**.

Rationality Leads to Morality

If we follow rationality consistently, we must act morally. This is because **morality is built into the very structure of rational decision-making**.

Key Questions

But does this mean that:

- **every rational person is automatically moral?**
- **Can a rational criminal exist?** |
- Is it rational for a company to pursue profit in a way that harms society?

Can a Rational Criminal Exist?

The Rationality of Immorality?

At first glance, it might seem **rational for someone to act immorally if it benefits them**. A criminal might steal money or commit fraud **to become wealthy**—after all, **isn't it rational to maximize one's own success?**

Kant's Response

However, Kant argues that this is **not true rationality**.
Why?



This is **a contradiction**—**true rationality demands universal consistency**, and crime cannot be justified this way

Rationality Requires Consistency

A rational person must follow principles that they would **accept if applied to everyone**.

If someone steals, they are implicitly saying, **"Stealing is acceptable."**

But if universally accepted, trust in property would collapse, and criminals would be vulnerable to theft too.



Self-Interest and True Rationality

The Problem with Narrow Self-Interest

Self-Interest at the Expense of Others is Irrational

Pursuing personal gain while disregarding moral laws is **self-defeating in the long run.**

Societal Consequences

If businesses, governments, or individuals **ignored morality, society would collapse into chaos**—which is **not in anyone's rational interest.**



Kant's Conclusion

A criminal, even if intelligent and calculating, **is not truly rational**—because their actions violate universal moral principles and ultimately **lead to contradiction and instability.**

Why Rationality Requires Morality?

Kant's argument follows a logical structure:

Rationality Requires Consistency

If a person is rational, they must be logically consistent.

Consistency Requires Universalizability

A rational person follows principles they would accept if everyone followed them.

Universalizability Requires Morality

If a rule is **not universally acceptable** (e.g., lying, stealing, harming others), it's **not a valid moral law**.

Rationality Leads to Morality

Therefore, if you are truly rational, you must act morally.

Immorality is Irrational

If you act immorally, you are acting irrationally.



Limitations of Universalizability

- Does Universalizability **Always Provide Acceptable Outcomes?**

A key challenge is whether **universalizability** is enough to **determine moral correctness**.

Just because I can universalize a maxim, does that make it good?

- Examining Test Cases**

Are there **tricky cases?**

Case Examples for Universalizability

Lying

- If everyone lies, communication would break down, trust would vanish—so **lying fails the universalizability test and is immoral.**

Robbing

If **robbing** were universalized, **property rights would collapse**, and society would descend into chaos—again, irrational and immoral.

Genocide

A genocidal leader might claim: "I believe certain people should be eliminated, and I would accept this rule universally."


Technically, **this maxim passes universalizability—but it is still deeply immoral.**

This suggests that **universalizability alone is not enough**—we also need to consider **human dignity and respect for life.**

Beyond Universalizability

Additional Criteria

Kant recognized that universalizability needs **supplementation with other principles.**



The Formula of Humanity

This is why Kant's second formulation of the **categorical imperative (treating persons as ends)** is crucial.

Combined Approach

Together, universalizability and **respect for humanity** provide **stronger moral guidance.**

While the principle of universalizability helps identify many immoral actions, it may fail to capture all moral intuitions on its own. A complete Kantian approach requires **both universalizability and respect for human dignity.**

Applying Kantian Ethics to AI Issues



AI and Privacy

How might Kant's principles guide the design of AI systems that collect and process personal data? **The principle of respecting humanity would require maintaining user autonomy and informed consent.**



Autonomous Vehicles

Should self-driving cars be programmed with **categorical rules (never harm a pedestrian) or utilitarian calculations (minimize overall harm)**? Kant would likely favor **strict rules respecting human dignity.**



Social Media Ethics

How can platforms respect user autonomy while preventing harm? The **universalizability** principle might suggest **policies that could be consistently applied to all users.**

Critiques of Kantian Ethics

1 Too Rigid

Critics argue that **Kant's absolute rules don't accommodate the complexity of moral life**. The famous "murderer at the door" scenario highlights this critique.

2 Conflicts Between Duties

What happens when categorical imperatives conflict? If I must never lie and always protect innocent life, what do I do when lying would save a life?



Critiques of Kantian Ethics

3 Difficulty in Application

Testing maxims for universalizability can be complicated in practice. **Different formulations of the same action might pass or fail the test.**

4 Neglect of Consequences

By focusing purely on **intentions** and **rules**, Kant's system might **permit actions with disastrous outcomes simply because they followed a good rule.**





Strengths of Kantian Ethics

1 Consistent Foundation

Kant provides a **systematic basis for morality grounded in reason** rather than fluctuating emotional responses or cultural relativism.

2 Universal Scope

The **categorical imperative** applies to all **rational beings**, providing a **framework for consistent moral judgment across different cultures and contexts**.



Strengths of Kantian Ethics

3 Respects Human Dignity

By emphasizing that people must be treated as **ends**, Kant's ethics provides **strong protection for human rights and dignity**.

4 Avoids Utilitarian Problems

Unlike consequentialism, Kantian ethics **prevents justifying harmful actions merely because they benefit the majority**.

Kantian Ethics and Contemporary Rights

Human Rights Foundation

Many **modern conceptions of human rights draw from Kant's idea** that all persons have inherent **dignity** that must be respected.

Consent and Autonomy

Contemporary emphasis on **informed consent** in medicine and research reflects Kant's respect for **autonomy** and treating persons as ends.

Global Ethics

International agreements on human rights often invoke **universal principles** that would be recognizable to Kant, even if not explicitly Kantian.



AI Dignity and Rights

1

Current Status

Today's AI **lack the rational autonomy** that Kant considered **necessary** for **dignity and moral status**.

2

Future Possibilities

If AI were to develop genuine moral reasoning and autonomy, **would Kant's framework grant them moral status?**

3

Ethical Implications

How might recognizing AI dignity (if possible) **change how we develop and deploy artificial intelligence?**

Kant's focus on rationality as the basis for dignity raises fascinating questions about the potential moral status of advanced AI systems. **As technology evolves, philosophers continue to debate whether artificial beings could ever possess the kind of rationality that confers moral worth.**

Balancing Deontology and Consequences



While Kant's ethics emphasizes **duty over consequences**, many contemporary ethicists seek a **middle ground that honors absolute principles while remaining sensitive to outcomes**. This **"threshold deontology"** might hold that **principles should generally be followed but could be overridden in extreme circumstances where consequences would be catastrophic**.

Do we want Kantian robots



Yes

- They will be consistent
- They will be impartial

No

- They may act on bad maxims
- Their maxims may be too rigid

David Ross: The Theory of Prima Facie Duties

Sir William David Ross (1877–1971)


Sir **William David Ross** (1877–1971) was a British philosopher who developed a **pluralistic deontological ethics**.

Prima Facie Duties

He introduced the idea of **prima facie duties**, which are **moral obligations that are binding unless they are overridden by a more important duty in a given situation**.

Unlike Immanuel Kant

Unlike Immanuel Kant, proposing that moral duties are absolute and universal, **Ross** argued that we often face **conflicting moral obligations** that require careful judgment.



David Ross: 7 Prima Facie Duties

Ross identified **7 key moral obligations** that we should generally follow. These duties guide our actions, though they are not always absolute.

1. Fidelity (Keeping Promises & Truthfulness)

We have a duty to **keep our promises** and **be honest** in our dealings with others.

Example: If you promise friends to help them move, you are morally obligated to do so unless a more pressing duty (e.g., an emergency) arises.

2. Reparation (Making Amends for Wrongs)

If we have **wronged someone**, we have an **obligation to make up for it**.

Example: If you damage someone's property, you should compensate or fix it.

3. Gratitude (Acknowledging and Returning Kindness)

When someone helps us, we should show **appreciation** and, if possible, **reciprocate the kindness**.

Example: If a colleague covers your shift, it would be moral to return the favor when they need help.



David Ross: 7 Prima Facie Duties

4. Non-Injury (**Non-Maleficence**: Avoiding Harm)

We must **not harm others**, whether **physically** or **psychologically**.

This duty is often considered more urgent than duties like beneficence (helping others).

Example: Even if honesty is important (fidelity), telling someone a truth that would cause unnecessary distress may not be the right action.

5. **Beneficence** (Promoting Well-being & Helping Others)

We have a duty to **help others** and contribute to their well-being.

Example: Donating to charity or assisting a stranger in need.

David Ross: 7 Prima Facie Duties

6. **Self-Improvement** (Developing Ourselves Ethically & Intellectually)

We should strive to **improve** our **health, wisdom, security, happiness, and well-being**.


Example: Taking time to study, exercise, or develop new skills is a moral obligation to oneself.

7. **Justice** (Fairness in Distributing Benefits and Burdens)

We should act **fairly** and **ensure benefits and responsibilities are distributed equitably**.

Example: Employers should treat workers fairly, ensuring equal pay for equal work.





Ross vs Kant: A Comparison

Aspect	David Ross (Prima Facie Duties)	Immanuel Kant (Deontological Ethics)
Nature of Duties	Moral duties are conditional (prima facie) and may conflict. Moral judgment is needed to determine priority.	Moral duties are absolute and universal , applying in all situations without exception.
Moral Decision-Making	Based on context and balancing competing duties .	Based on strict rules (categorical imperative).
Example: Lying to Protect a Friend	If honesty (fidelity) conflicts with preventing harm (non-injury), the duty to avoid harm may override truthfulness.	Lying is always wrong , no matter the consequences.
Flexibility	More flexible , as duties can be overridden based on circumstances.	Rigid and absolute —moral laws must always be followed.



Key Takeaways

Ross acknowledges moral complexity

In real life, duties often clash, and we must **weigh** them carefully.

Kant insists on universal moral laws

Moral rules are **absolute** and must be followed **regardless of the consequences**.

Practical differences

Ross's approach is **more adaptable** to everyday moral dilemmas, whereas **Kant's framework is more rigid** but offers clear guidance.



Defeasibility of duties

The concept of defeasibility

Does it make sense to view duties as being defeasible (i.e., decisions can be revised if new info arises)?

Contextual reasoning

We start with a general moral rule (e.g., "We should tell the truth"). However, **new context** may provide a compelling reason to override that duty. The conclusion is **defeasible**—it can change based on further evidence or competing moral principles.

Practical applications

Can we apply defeasible reasoning to reason with duties?

Should an AI system admit exceptions to duties, or should it always ask humans?



Examples of Defeasible Reasoning in Prima Facie Duties

Fidelity vs. Non-Injury

Suppose a doctor knows that a patient has a **terminal illness**. Should they disclose the full truth if it might cause severe emotional distress or even suicide?

Normally, **fidelity** (truthfulness) demands that the doctor be honest.

However, **non-injury** (avoiding harm) suggests that **in some cases**, the truth should be disclosed gradually or with caution.



Examples of Defeasible Reasoning in Prima Facie Duties

Reparation vs. Justice

Imagine you accidentally **break** a friend's valuable item. You owe **reparation** (compensating for the damage).

However, if you are **financially struggling** and paying full compensation would leave you **unable to afford necessities**, **justice** might require a fairer solution (e.g., a **delayed or partial payment**).



AI examples

Option 1: AI Systems Allowing Exceptions (Defeasible AI Ethics)


AI could be programmed to use **defeasible reasoning** to recognize **context-dependent exceptions** to moral rules.

Example:

- A self-driving car has a **duty to obey traffic laws**.
- However, if an **emergency situation** arises (e.g., an ambulance behind it, needing to yield), the AI could recognize that **justice (saving lives) overrides strict rule-following**.

Risks:

- AI might make **incorrect ethical judgments**, leading to unintended consequences.
- Moral complexity might make AI decisions **unpredictable or inconsistent**.



AI examples

Option 2: AI Always Asking Humans

The AI system could be designed to **always defer ethical dilemmas to human oversight** rather than making independent decisions.

Example:

- A medical AI detecting a **life-threatening but ambiguous case** could notify a human doctor instead of deciding on treatment autonomously.

Risks:


- Slower decision-making in **time-sensitive situations** (e.g., autonomous vehicles, emergency medicine).
- **Human bias** or **fatigue** could lead to inconsistent decisions.

AI examples

A **Balanced Approach?** Hybrid AI with Human-in-the-Loop


A more **practical approach** is a **hybrid AI model**:

- AI **uses defeasible reasoning** to handle **routine ethical exceptions** (like an autonomous vehicle yielding to emergency responders).
- For **high-stakes dilemmas**, the AI **consults a human** (e.g., medical AI recommending but not autonomously executing treatment plans).



Nietzsche: a critique of ethics

Friedrich Nietzsche was a **radical critic of traditional ethics**, particularly the moral systems derived from Christianity and Enlightenment philosophy. He believed that **conventional morality** – which values humility, equality, and self-sacrifice – is **a tool of the weak to restrain the strong**. Instead, he proposed an alternative vision centered on **the Übermensch (Overman/Superior Human)**, a figure who **creates values rather than obeying them**.



Nietzsche: a critique of ethics

Key Concepts in Nietzsche's Critique of Ethics

1. The **Übermensch**: **Beyond Good and Evil**

The **Superior Human** is an individual who **rises above traditional moral values** of good and evil.

Traditional morality, according to Nietzsche, is a "**morality of the herd**"—a system imposed by the weak to control the strong.

Instead of conforming to societal norms, the **Übermensch defines their own values** based on **strength, creativity, and personal excellence**.

Example:

A great **artist, scientist, or leader** does not ask society for approval; they **create their own path** and reshape the world according to their vision.



Nietzsche: a critique of ethics

Key Concepts in Nietzsche's Critique of Ethics

2. Duties Only Toward Equals

Nietzsche **rejects** the idea of **universal moral obligations**. Instead, he claims:

- One **only** has **duties toward equals**, meaning those who are also strong, independent, and capable.
- Toward **beings of lower rank**, one acts as one sees fit—**not out of obligation, but as an expression of personal will**.

This is a **direct challenge** to moral theories like Kant's **universal moral law** or Ross's **prima facie duties**, which apply to all individuals equally.

Example:

A powerful leader or thinker should not be bound by the same rules as an ordinary person. Their **responsibility** is to **their own potential**, not to the demands of the weak.



Nietzsche: a critique of ethics

Key Concepts in Nietzsche's Critique of Ethics

3. Value Creation: The Superior Human Determines Right and Wrong

Unlike philosophers believing in **objective moral values** Nietzsche claims that **values do not exist independently**—they are **created**.

The **Übermensch** does not **passively accept morality** but **actively determines what is valuable**.

Example:

A visionary entrepreneur, like Elon Musk does not follow pre-established business ethics but **redefines innovation, competition, and ambition** according to their own vision.

Instead of asking, "What is right?" or "What is wrong?" Nietzsche's superior individual asks, **"What do I will?"**



Nietzsche: a critique of ethics

Key Concepts in Nietzsche's Critique of Ethics

4. No Need for Justification: The Will to Power

The superior human **does not need external validation** (from God, society, or philosophy) for their moral choices.

The only criterion of **wrongness** is:

- "That which is harmful to me is harmful as such."

This aligns with **Nietzsche's concept of the Will to Power**—the fundamental drive for growth, dominance, and self-overcoming.

Example:

A revolutionary thinker like **Galileo did not wait for society's approval** to challenge existing beliefs. He **pushed forward** despite opposition because he valued **truth and discovery above traditional dogma**.

Nietzsche vs. Other Ethical Theories



Ethical Theory	Core Idea	Nietzsche's Critique
Kant (Deontology)	Moral rules are absolute and apply to everyone.	The strong should not be bound by the same rules as the weak.
Utilitarianism	Morality is about maximizing happiness for all .	Morality should not be dictated by majority well-being but by individual will .
Ross (Prima Facie Duties)	Duties are flexible and based on context.	Duties are meaningless unless self-imposed by the superior individual .

Implications & Controversies of Nietzsche's Ethics

✓ Empowerment & Individual Greatness

- Encourages **self-creation** and **personal growth** rather than conformity.
- Challenges people to **rise above societal norms** and redefine success on their own terms.

✗ Risk of Elitism & Moral Relativism

- Could justify **oppression** if the powerful claim they owe nothing to the weak.
- If each person creates their own values, is anything truly wrong?



FRIEDRICH NIEZSCHE

Nietzsche: a critique of ethics

Nietzsche's philosophy challenges us to question morality itself. **Is morality a universal truth, or just a construct to maintain order? Should we strive for personal excellence rather than moral obedience?** His ideas continue to inspire and provoke debate in modern philosophy, leadership, and ethics.



Contractarianism

Foundations of Justice and Morality

Between Political and Moral Theories

Social contract theories

In Political Theory

A societal arrangement is **just if** it has been (or would have been) **accepted by free and rational people**.

Key Idea:

A just society is not determined by divine command, tradition, or brute force but **by what free and rational individuals would agree to if they were designing society from a fair standpoint**.

This idea is central to **contractarian and contractualist traditions**, seen in thinkers like Thomas **Hobbes**, John **Locke**, Jean-Jacques **Rousseau**, and John **Rawls**.

Example:

Imagine designing a society from scratch. If rational people would agree to a particular set of laws (e.g., equal voting rights, property protections), then that arrangement is considered just.

Contrast in Thinkers:

- **Hobbes**: People give up freedoms to a strong authority (Leviathan) to ensure security.
- **Locke**: Governments must protect life, liberty, and property, and they derive power from the consent of the governed.
- **Rousseau**: The general will represents collective freedom, where laws are expressions of the people's will.
- **Rawls**: Justice is what rational people would choose behind a veil of ignorance (without knowing their own social position).

Social contract theories

In Moral Theory

Actions are morally right just because they are **permitted by rules that free, equal, and rational people would agree to live by**, on the condition that others obey these rules as well (Shafer-Landau).

Key Idea:

Moral rules **are not absolute truths** but **agreements that rational individuals would make** to create a fair and cooperative society.

If a moral principle is something that **everyone would reasonably agree to**, it is considered **justified**.

Examples:

- **Lying:** If everyone lies, trust collapses. Since rational individuals prefer a society where people tell the truth, **truthfulness** is a **moral duty** under the social contract.
- **Stealing:** People would **not** agree to a society where theft is allowed, so prohibiting theft is a **socially agreed moral rule**.

State of nature and social contract



The **state of nature** is a **hypothetical condition** where there is **no government, no laws, and no organized society**.

Classic philosophers who discussed the state of nature:

- **Thomas Hobbes:** Life in the state of nature is "solitary, poor, nasty, brutish, and short" because individuals act purely out of **self-interest**.
- **John Locke:** The state of nature is not as chaotic as Hobbes suggested, but still **lacks a secure system** of justice and property protection.
- **Jean-Jacques Rousseau:** Humans in the state of nature were peaceful, but **society introduced inequality and corruption**.

Key Question:

- How do we get out of the state of nature?
- Through **social contracts**, where individuals agree to rules that ensure **mutual benefits and security**.

PRISONER'S DILEMMA

State of nature and social contract

The matrix in the slide represents a classic Prisoner's Dilemma scenario, which helps explain why cooperation is difficult without a binding agreement.

A/B	Cooperate	Defect
Cooperate	-1,-1 (Mutual benefit)	-5,0 (A is exploited)
Defect	0,-5 (A exploits B)	-3,-3 (Mutual loss)

Understanding the Payoff Matrix

- 1 Best collective outcome: (Cooperate, Cooperate) = (-1,-1) → Both benefit.
- 2 Best selfish move: Defecting while the other cooperates (-5,0) or (0,-5).
- 3 Worst scenario: If both defect (-3,-3), neither benefits.

Splakket is yames
Ofcerient oudotyon!

State of nature and social contract



What does this mean?

- Without a **social contract**, **rational self-interest** leads individuals to defect, resulting in **mutual loss**.
- With a **binding agreement (a social contract)**, individuals can **commit to cooperation**, leading to **better outcomes**.

What Agreements Are OK?

- **Agreements** that encourage **trust and cooperation** while **detering defection**.
- A legitimate agreement must be **fair** and **mutually beneficial**, ensuring that all individuals have incentives to follow it.

State of nature and social contract

-  **A just agreement**
Fair, mutually beneficial, and enforceable
-  **Social contract**
Rules that make cooperation rational
-  **Prisoner's Dilemma**
Rational self-interest prevents cooperation
-  **State of nature**
Leads to distrust and conflict

Theory	Core Idea	Social Contract Perspective
Kant (Deontology)	Moral duties are universal and absolute.	Morality depends on agreements rather than absolute laws.
Utilitarianism	Right actions maximize happiness for all.	Morality is about fair rules that rational people accept , not just consequences.
Nietzsche (Critique of Ethics)	The strong create their own values.	Morality is collective rather than determined by individuals.

Strengths & Criticisms of Social Contract Theory

✓ Strengths:

- **Justifies political authority** in a way that respects **individual autonomy**.
- **Explains moral obligations** without appealing to religion or intrinsic moral laws.
- Provides a **fair way** to structure society based on **rational agreement**.

✗ Criticisms:

- **Who counts as a rational person?** Historically, some groups (**women, minorities**) were excluded.
- **Hypothetical consent:** Just because **rational people *would*** agree does not mean **real people *have*** agreed.
- **What about obligations to non-contracting beings?** (e.g., animals, future generations).

John Rawls (1921-2002)

Rawls' Theory of Justice

John Rawls was a 20th-century philosopher who proposed a **justice-based social contract** that prioritizes **fairness** and **equality**.

His key work, ***A Theory of Justice (1971)***, presents a framework for **how societies should structure laws and institutions** to be just.





John Rawls



Problem: How Do We Ensure a Fair Social Contract?



The veil of ignorance

The **veil of ignorance** as a method to ensure **fairness** in social contracts.



Choosing principles


People should **choose** **under a veil of ignorance**, **without knowing their gender, social position, interests, talents, wealth, race, etc.**

John Rawls and the Veil of Ignorance

A fair society is one that rational individuals would choose if they did not know their place in it.

How Does the Veil of Ignorance Work?

- Imagine you are **designing society's rules** but **without knowing your own identity** (gender, wealth, race, talents, or social class).
- Would you agree to extreme inequality if you might end up poor?
- Would you allow discrimination if you might belong to the discriminated group?



💡 **Rawls' claim:** Under the **veil of ignorance**, rational people would choose **just and fair principles**, as **they wouldn't risk ending up in a disadvantaged position.**

John Rawls' Two Principles of Justice

1. Equal Basic Liberties

Everyone should have **the same fundamental freedoms**

These liberties include:

- ✓ Freedom of **conscience** (religion, speech, thought).
- ✓ Freedom of **association** (political and social groups).
- ✓ Freedom of **speech** and **personal liberty**.
- ✓ Right to **vote and participate in democracy**.

Basic **liberties** must be **equal for all** and **cannot be sacrificed for economic or social advantages**.



John Rawls' Two Principles of Justice

2. Fair Equality of Opportunity & The Difference Principle

"Social and economic inequalities are to satisfy 2 conditions..."

Condition 1: Fair Equality of Opportunity

Opportunities for positions of power, influence, and wealth must be genuinely **open to everyone**, not just those born into privilege.

Example: If education is only accessible to the wealthy, **true equality of opportunity does not exist**. A **just** society ensures that **everyone has a fair chance** regardless of their background.

Condition 2: The Difference Principle

Inequalities in wealth and power are **acceptable only if** they **benefit the least advantaged** members of society.

Example: A CEO making millions is **only justifiable** if their success also improves conditions for workers and the poor. **Progressive taxation** (where the rich pay more to support social programs) aligns with this principle.

John Rawls and the Veil of Ignorance



Modern Policies

Welfare, affirmative action, and public healthcare systems apply Rawls' ideas to protect the disadvantaged.



AI & Ethics

Can AI systems make fair decisions without bias?
Should AI policies be designed as if we were behind a "veil of ignorance"?



Global Justice

Should richer countries **redistribute wealth** to poorer ones under the **difference principle**?



Rawls vs. Other Theories

Theory	How Society Should Be Structured	Fairness Test?
Hobbes (Contractarianism)	People obey a strong government to escape chaos.	No concern for fairness, just security.
Locke (Natural Rights)	Society protects life, liberty, and property.	Concern for rights, but assumes economic inequality.
Utilitarianism (Greatest Happiness Principle)	The goal is maximizing happiness, even if some suffer.	Could justify unfair distributions if they maximize total welfare.
Rawls (Justice as Fairness)	Society must be fair, ensuring basic liberties and helping the least advantaged.	Yes! The veil of ignorance ensures fairness.

Final Takeaway

-  **Basic liberties are non-negotiable**—everyone must have equal fundamental rights.
-  **Opportunities must be fair**—success should not depend on birth status.
-  **Inequalities are only justifiable if they benefit the least advantaged.**

AI in a just society according to Rawls


Does AI Deployment Fit Rawls' Requirements?

AI is transforming **work, healthcare, governance, and decision-making**, but **does it uphold the principles of justice?**

To be just under **Rawls' framework**, AI must:

- **✓ Respect basic liberties** (freedom of speech, privacy, due process).
- **✓ Ensure fair equality of opportunity** (no AI-driven discrimination).
- **✓ Benefit the least advantaged** (aligning with the difference principle).

If AI **reinforces inequality** or limits individual freedoms, it **violates** Rawls' vision of justice.



AI in a just society according to Rawls

When May AI Conflict with Basic Liberties?

Rawls' First Principle: **Each person must have equal fundamental rights and freedoms.**

Possible AI Violations of Basic Liberties:

✗ **Privacy Violations:**

AI-driven surveillance (e.g., facial recognition) may **limit individual privacy and freedom of movement.**

Automated decision-making can monitor and control behavior, leading to **potential government overreach.**

Freedom Issues

✗ **Freedom of Expression & Political Rights:**

Social media algorithms may **manipulate public discourse, distorting democracy.**

AI bias in content moderation can lead to censorship and restrictions on speech.





AI in a just society according to Rawls

Law Enforcement Concerns

✗ Bias in Law Enforcement & Criminal Justice:

Predictive policing algorithms have been shown to reinforce racial and socioeconomic biases, undermining equal treatment under the law.

💡 Rawlsian Perspective: AI should be regulated to protect personal freedoms and prevent discrimination.

AI in a just society according to Rawls

When May AI Conflict with Fair Equality of Opportunity?

*Rawls' Second Principle
(Condition 1): Access to jobs and social positions should not be limited by arbitrary factors like race, gender, or wealth.*

How AI Can Undermine Equal Opportunity:

✗ Algorithmic Bias in Hiring & Lending:

- AI-driven hiring tools have been found to **discriminate against women and minorities** based on biased training data.
- AI in financial systems may **deny loans based on zip codes or income history**, reinforcing **economic inequality**.





AI in a just society according to Rawls

Education and Employment Concerns

✗ Education & Job Market Displacement:

- AI-driven automation disproportionately impacts **low-income workers**, eliminating jobs **without adequate retraining programs**.
- If AI **exacerbates the digital divide**, only the privileged will **benefit from AI-driven education and career opportunities**.

Rawlsian Perspective:

- AI must be designed to **correct, not reinforce, social inequalities**.
- Equal opportunity means **ensuring access to AI-driven education, fair hiring practices, and retraining for displaced workers**.

AI in a just society according to Rawls

When May AI Conflict with the Difference Principle?

Rawls' Second Principle (Condition 2): Economic inequalities are only just if they benefit the least advantaged.

How AI Can Worsen Inequality:

✗ **Wealth Concentration in AI Corporations:**

- AI's economic benefits largely **favor tech giants and elites**, widening the rich-poor gap.
- If AI-driven automation **replaces jobs without compensation for workers**, the wealthiest **benefit while the working class suffers**.



AI in a just society according to Rawls



Healthcare Concerns

✗ Healthcare AI Access Disparity:

- **AI-powered** healthcare tools (e.g., diagnostics, robotic surgery) **may only be available to the wealthy**, increasing **health inequalities**.
- Unequal access to AI-driven treatments would violate the **difference principle**.

💡 Rawlsian Perspective:

- AI-driven economic growth must **actively benefit disadvantaged groups** through **redistribution policies**, **affordable AI-based healthcare**, and **public AI education**.
- If AI creates wealth, it **must uplift the poorest**, **not just benefit corporations**.





Juergen Habermas: Discourse Ethics

Juergen Habermas (born 1929) is a **German philosopher** known for his work on **ethics, democracy, and communication theory**. His **discourse ethics** proposes that **moral norms and rules** are only valid if they are **agreed upon through rational and free discourse** among all affected individuals

Juergen Habermas: Discourse Ethics

Core Idea: Justification Through Discourse

*A rule of **action or choice** is **justified**, and thus valid, only if **all those affected by it could accept it in a reasonable discourse**.*

Instead of grounding morality in **divine authority** (religion), **rational duty** (Kant), or **consequences** (Utilitarianism), Habermas argues that **justification emerges through rational discourse**.

A norm is valid only if **all affected individuals could accept it in an open discussion** without coercion or manipulation.



Example:

A government **imposing a policy** (e.g., raising taxes for social programs) must **justify it through public discourse**, ensuring that **all affected individuals** have a say.

Juergen Habermas: Discourse Ethics

- Valid Norms & the "Ideal Speech Situation"
A norm is valid when the foreseeable consequences and side effects of its general observance for the interests and value orientations of each individual could be jointly accepted by all concerned without coercion.
- ✓ Free from coercion (no threats or manipulation)
- ✓ Based on reason and evidence (not power, emotions, or social status)
- 💡 Example:
In a truly democratic society, laws and policies should be based on public deliberation, scientific evidence, and fair reasoning, not political power struggles.
- The "Ideal Speech Situation" (ISS):
For discourse ethics to work, communication must be:
- ✓ Open to all perspectives (all affected individuals can participate)
- ✓ Aimed at consensus (seeking mutual understanding, not winning an argument)
- 🚫 Problem: Real-world discussions are rarely free from bias, misinformation, or unequal power dynamics.

Kant's alternatives
of Morality.



Juergen Habermas Discourse Ethics vs Other Moral Theories

Theory	Basis of Morality	Habermas' Critique
Kant (Deontology)	Morality is based on universal moral duties (Categorical Imperative).	Too rigid—does not consider dialogue or social context .
Utilitarianism	Morality is about maximizing happiness.	Ignores the right to justification —people must be able to debate rules , not just accept what is "best for most".
Rawls (Justice as Fairness)	A just society is based on what rational people would choose behind the veil of ignorance.	Habermas agrees that fairness matters but insists that fair rules must emerge from actual discourse .



Juergen Habermas: Discourse Ethics

Why Does Discourse Ethics Matter Today?

✓ **Democratic Decision-Making:** Encourages **open** debates, **free press**, and **public participation** in policy-making.

✓ **Social Media & Public Discourse:** In a digital age where algorithms influence discussion, we must **protect spaces** for **rational, inclusive discourse**.

✓ **AI & Ethics:** Ethical AI should be **transparent**, **explainable**, and **open to public scrutiny** to ensure **fair** decision-making.

🚩 Challenges:

Power imbalances (wealthy corporations dominate discourse).

Misinformation & manipulation (fake news, political propaganda).

Limited participation (not everyone has access to free speech or education).

Habermas and AI



Equal Participation

All affected by AI **must have a voice**



Rational Discussion

Based on **reason** and **evidence**



Inclusive Framework

All stakeholders included

Would we all agree if we engaged in an impartial discussion on how to use AI?

According to Habermas, a just system requires **rational, inclusive, and coercion-free discourse**. If we applied this to AI, **a fair discussion about AI's role in society would mean:**




Equal Participation:

All individuals affected by AI (citizens, workers, policymakers, developers) **must have a voice** in shaping AI policies.

AI decisions should **not** be left **only** to **tech companies or governments**.

AI System not, int a. Trasplart
Ethical Decision -making Eigrntes,
Proces s, exthital -naking aligimims.



Habermas and AI

✓ Transparency & Justification:

AI should be **explainable**—its **decisions** must be **understandable** and **justifiable** in public discourse.

No "black box" decision-making that excludes public scrutiny.

✓ Protection Against Bias & Coercion:

AI should **not manipulate opinions** (e.g., through biased recommendation algorithms).

AI must avoid reinforcing **existing social inequalities** (e.g., biased hiring tools).

✓ AI Must Serve the Public Good:

AI should be designed to **benefit all**, not just corporations or powerful governments.

AI use must be **debated democratically**, ensuring its impact aligns with shared ethical values.

Habermas and AI



✓ Facilitator:

AI could help **structure debates**, ensuring fairness (e.g., moderating online discussions, fact-checking claims).

The Challenge:

AI can process **data and arguments**, but **can it genuinely reason** and **engage in free discourse** like humans?

Can it understand fairness beyond statistics and logic?



What AI Lacks?

Personal moral intuition? Does AI have **lived experiences**, emotions, or moral agency?

True impartiality ? AI systems are trained on human data, meaning **they inherit biases unless actively corrected?**



✓ Impartial Advisor:

AI could analyze ethical arguments and highlight **logically (in)consistent** positions.



✓ Simulated Perspectives:


AI could model different ethical frameworks (utilitarianism, Kantian ethics, Rawlsian justice) to **provide multiple viewpoints**.



Virtue ethics




Francesca Lagioia



Virtue ethics

Virtue ethics is an **ethical framework** that focuses on **character** rather than rules or consequences.

It originates from **Aristotle's philosophy**, emphasizing the **development of virtues to achieve a flourishing life** (eudaimonia).



Virtue ethics

1. Ethics Should Not Focus on Norms or Consequences

An act is morally right just because it is one that a virtuous person, acting in character, would do in that situation.

It asks:

💡 What would a virtuous person do in this situation?

Unlike deontology (Kant) or utilitarianism (Mill), **virtue ethics** does **not** focus on:

✓ **Fixed moral rules** (like Kant's categorical imperative).

✓ **Maximizing consequences** (like utilitarianism).

✓ **Example:**

A doctor treating a patient does not only follow rules (deontology) or calculate consequences (utilitarianism) but acts with **compassion, honesty, and courage** to make the right decision.



Virtue ethics



2. Ethics is a Complex Matter

Since there are many virtues, the right act is the one resulting from the mix of the relevant virtues: honesty, loyalty, courage, impartiality, wisdom, fidelity, generosity, compassion, etc.




Moral actions depend on **context** and the **balance of virtues**.

Example:

- Telling the truth (**honesty**) is good, **but** if the truth would cause unnecessary harm, **compassion** might override honesty.
- A leader must **balance courage** with **wisdom**—reckless bravery is not a virtue.



No single virtue is always correct—**wisdom helps balance competing virtues**.



Virtue ethics



3. Ethics Cannot Be Learned Through a Set of Rules

Its application requires practical wisdom (phronesis).

Example:

A young person does not become **honest** just by memorizing moral principles—they **learn by practice**, observing honest individuals, and developing **good habits**.



Virtue is learned through **experience**, **habituation**, and **role models**, not just by following rules.



Moral education and **practice** help individuals **develop good character**.



Virtue ethics and Wisdom

For Aristotle, **practical wisdom** is an **intellectual virtue** (not a moral virtue) that enables a person to: correctly assess concrete circumstances

- choose the right means toward a good end
- find the golden mean between excess and deficiency
- integrate emotions, reason, and character

It is not mere “intelligence” nor simply technical skill. It is **the capacity to perceive a situation morally.**

Aristotle distinguishes between:

- **Episteme** → scientific knowledge (universal and necessary)
- **Techne** → technical skill (knowing how to make or produce something)
- **Phronesis** → practical wisdom (knowing how to live well)



Virtue ethics and Wisdom

Practical wisdom is **not**:

- a **technique** (it is not “how to manipulate effectively”)
- nor merely **theoretical knowledge**

It is a capacity for situated moral judgment.

- Virtues (courage, justice, temperance...) do **not** operate **automatically**.
- They require **discernment** and **moral sensitivity** to be applied properly.

Example:

- Telling the truth is a virtue.
- But telling the truth in a brutal way can become cruelty.
- **Practical wisdom allows us to understand how, when, and how much truth to tell.**
- Without phronesis, virtues become rigid or even dangerous.

How Virtue Ethics Compares to Other Ethical Theories

Theory	Core Idea	Virtue Ethics Critique
Deontology (Kant)	Moral duties are universal and must always be followed.	Too rigid—does not consider context or moral character .
Utilitarianism (Mill)	The right action maximizes happiness .	Too focused on outcomes , ignoring personal integrity .
Virtue Ethics (Aristotle)	Ethics is about developing good character and making decisions as a virtuous person would.	More flexible , but lacks clear-cut rules for action .



Issues in Virtue Ethics

Virtue ethics offers a **flexible and character-based approach** to morality, but it also **raises** several **philosophical challenges**.

.

Who decides **what counts as a virtue**? Different cultures and societies value **different traits**.

1. How Do We Know What is Virtuous and What is Not?



The Problem:

- Unlike deontological ethics (fixed rules) or utilitarianism (clear outcomes), **virtue ethics lacks a strict formula** for determining right and wrong.

Issues in Virtue Ethics

Example:

- **Ancient Greece** valued **courage and honor**, while **modern societies** may prioritize **compassion and inclusivity**.
- Is **ambition** a virtue (driving progress) or a vice (causing selfishness)?

Possible Solution:

Look at **historical and cultural perspectives** but also engage in **philosophical reflection** to **adapt virtues to modern contexts**.



Issues in Virtue Ethics

2. How Can We Extract Precise Moral Guidance from Virtuous Examples?

🤔 The Problem:

Virtue ethics relies on **role models** (e.g., Aristotle's "virtuous person"), but how do we ensure **consistency**?

Traditions **evolve over time**—should we rely on **historical figures**, or do we need **new role models**?

💡 Example:

If we look to **Socrates, Gandhi, or Martin Luther King Jr.**, their virtues seem clear.

But what about **leaders with both virtues and flaws** (e.g., Churchill, Steve Jobs)?

Can we separate **good virtues** from **bad traits** in historical figures?

🔍 Possible Solution:

Use **phronesis (practical wisdom)** to interpret **contextually relevant virtues** rather than blindly following tradition.

Issues in Virtue Ethics

3. What if Virtues Are in Conflict?

🤔 **The Problem:** Since virtue ethics **balances multiple virtues**, what happens when two virtues **conflict**?

💡 Example:

A journalist must balance **honesty** (revealing the truth) and **compassion** (not harming someone's reputation).

🔍 Possible Solution:

Practical wisdom (phronesis) helps decide which virtue is **most relevant** in a specific case.

There is **no clear rule** on which virtue should take priority.

A doctor faces a conflict between **justice** (fair treatment for all) and **care** (giving extra attention to a struggling patient).

Consider the **context and long-term consequences** of prioritizing one virtue over another.

Issues in Virtue Ethics

4. What Are the Paradigms of Virtue?

🤔 The Problem:

Are virtues **universal** (the same across cultures) or **relative** (shaped by history and society)?

Should we rely on **ancient virtue traditions (Aristotle, Confucius)** or develop **new ethical paradigms**?

💡 Example:

Some virtues are **consistent across cultures** (e.g., **courage, honesty, generosity**).

Others **vary**:

Western ethics emphasizes **individual autonomy**.

Eastern ethics values **harmony and community**.

🔍 Possible Solution:

Engage in **global ethical discourse** to identify **core virtues** while respecting **cultural diversity**.

Virtue Ethics: Takeaway

1

Context-Sensitive

Virtue ethics is **context-sensitive** but **lacks clear-cut moral rules**.

2

Practical Wisdom


We need **practical wisdom (phronesis)** to **navigate moral complexity**.

3

Adaptation

Virtue **traditions** must be **critically examined and adapted** for modern challenges.

Would you like **concrete applications of virtue ethics in leadership, AI, or public policy?** 🚀



AI and virtue ethics

Virtue ethics emphasizes **character development** and **moral wisdom** rather than fixed rules or consequences.

Applying it to **AI ethics** raises important **questions** about the **moral responsibilities of developers and AI systems**



AI and virtue ethics

1. Should AI Developers Be Virtuous? What Traits Should They Cultivate?

💡 **Developers** of AI systems shape technology that impacts society. Should they **embody virtues in their work?**

✓ Key Virtues for AI Developers:

- **Honesty:** Avoid **misleading claims about AI's capabilities.**
- **Justice:** Ensure **fairness**, in AI decision-making.
- **Responsibility:** Be accountable for the ethical implications of AI.
- **Wisdom (Phronesis):** Use good judgment in balancing risks and benefits.
- **Compassion:** Consider the well-being of those affected by AI.

✓ **Example:** Developers working on AI for hiring should **actively test for bias** to prevent discrimination.

AI and virtue ethics

2. Should AI Systems Be Virtuous?

💡 If AI influences **healthcare, hiring, law enforcement, and social interactions**, should it be programmed to **act virtuously**?

Risk of manipulation

✗ AI with "virtuous" behavior could be used to **control public opinion**.

✓ Example:

- Should **AI-powered hiring tools** **prioritize fairness over efficiency**?
- Should **autonomous weapons** be programmed with **compassion and restraint**?

Challenges

✗ **Can AI truly develop character?**

Virtues are based on **moral understanding**, not just rules or patterns.

Who defines virtue?

✗ Different societies prioritize different virtues.

AI and virtue ethics

3. How Can Virtues Be Learned?

💡 In humans, virtues are developed through **practice, experience, and moral reasoning**.

Value alignment

Ensuring AI is aligned with **human moral standards**.

Reinforcement learning

Rewarding "virtuous" decisions in AI's training.

Mimicking moral role models

Training AI on **ethical human behavior**.

🚩 Limitations:

- AI **lacks true moral understanding**—it can only **simulate** virtue.
- Ethics is **context-sensitive**, but AI relies on **fixed patterns**.

✓ **Example:** An AI **trained on compassionate healthcare decisions** might **prioritize patient well-being**, but would it truly understand **why** compassion matters and adapt its behavior if needed?

AI and virtue ethics

4. Can Training AI Lead to Virtuous Behavior?

💡 If **virtue requires practical wisdom**, can AI develop **wisdom through training**?

✅ **Yes, to some extent:**

- AI can be trained to **detect and correct biases**.
- AI can **simulate ethical reasoning** based on **historical data**.

❌ **No, fully developing virtue requires:**

- **Conscious moral reflection**, which AI lacks.
- **Adapting to new ethical challenges**, rather than relying on pre-set patterns.

✓ **Example:**

- A **self-driving car** can be trained to prioritize safety (prudence), but can it make **moral trade-offs** in ambiguous situations?
- A **chatbot** can be trained to show empathy, but does it **genuinely care** or just **simulate concern**?




AI and virtue ethics

Takeaway

- ☐ **✓ Developers should embody virtues**—they are responsible for AI's societal impact.
- ☐ **✓ AI can be trained to act virtuously**, but **it cannot truly develop moral character**.
- ☐ **✓ Virtue ethics challenges us to ask:** Should AI **imitate** virtue, or should it be left to **human moral judgment**?

Readings



- Shafer-Landau, R. (2018). The Fundamentals of Ethics. Oxford University Press.
- Singer, P. (2021). Ethics. In Encyclopedia Britannica: <https://www.britannica.com/topic/ethics-philosophy>